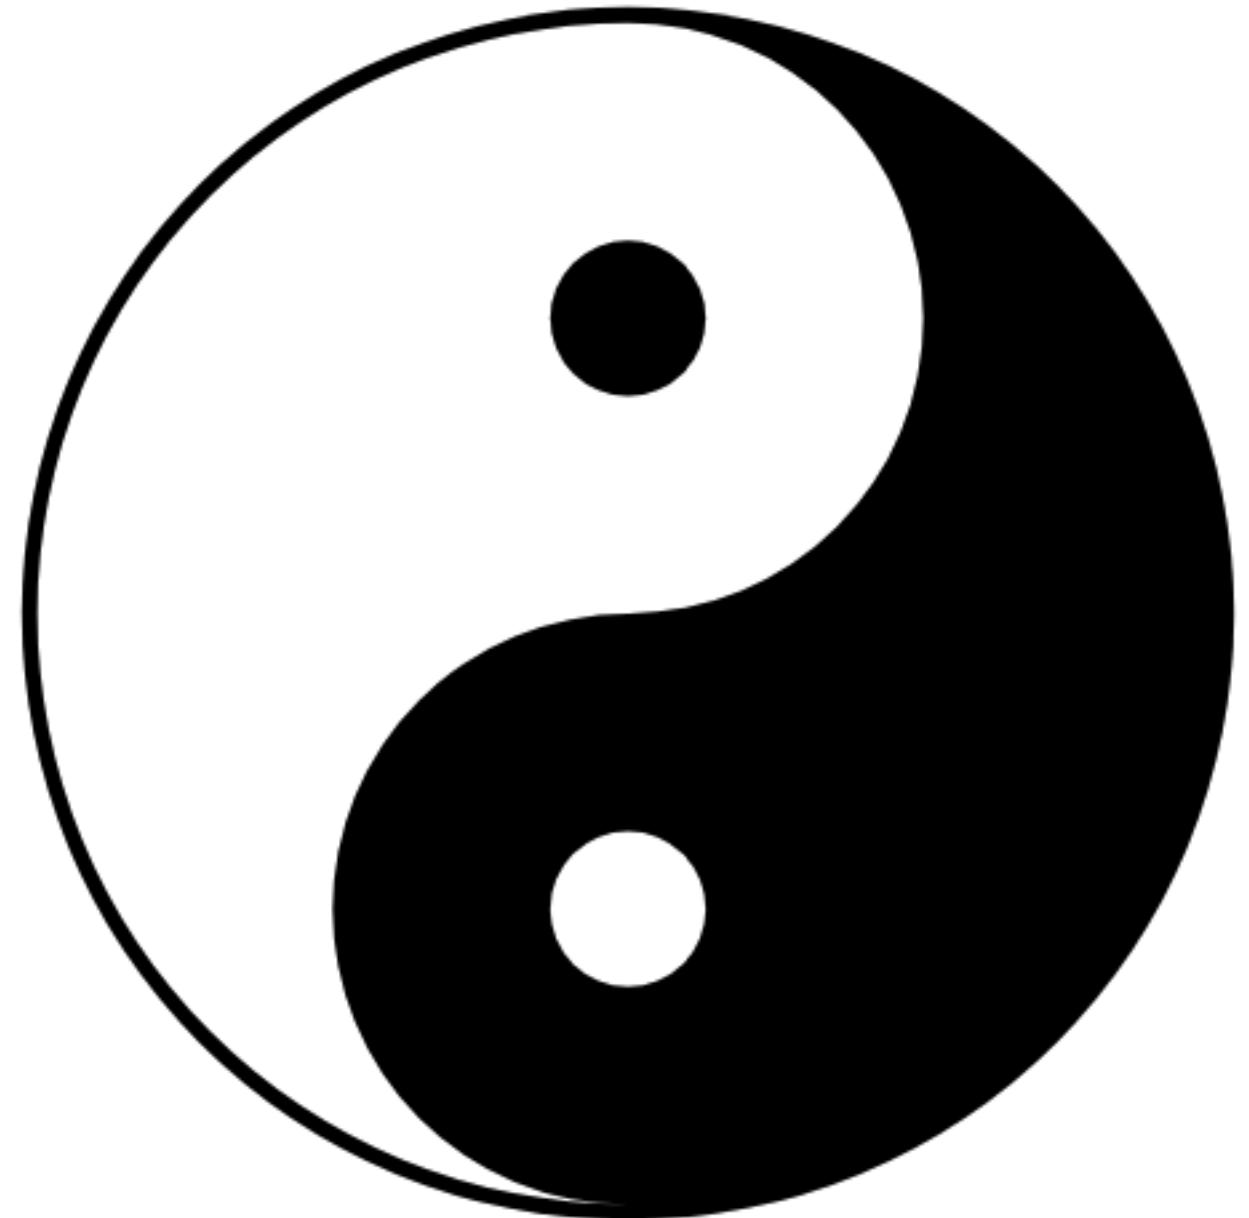


Yin Yang Ranch: A Distributed Computer Vision System with Raspberry Pi's and Macs



Jeff Bass

Building a Small Permaculture Farm in Suburbia

- Permaculture is a collection of practices to make farming more sustainable by building living soil, emulating old growth forests and recycling nutrients and water
- I am turning my 2 acre suburban lot into a small permaculture farm called Yin Yang Ranch
- It is an evolving science project and a demonstration garden; lots of shared learnings with others
- Growing a plant polyculture: figs, pears, pomegranates, plums, grapes, avocados, oranges, mulberries, blackberries...
- ... and a bumper crop of Raspberry Pi's.
- Using Computer Vision as a tool to help manage the farm



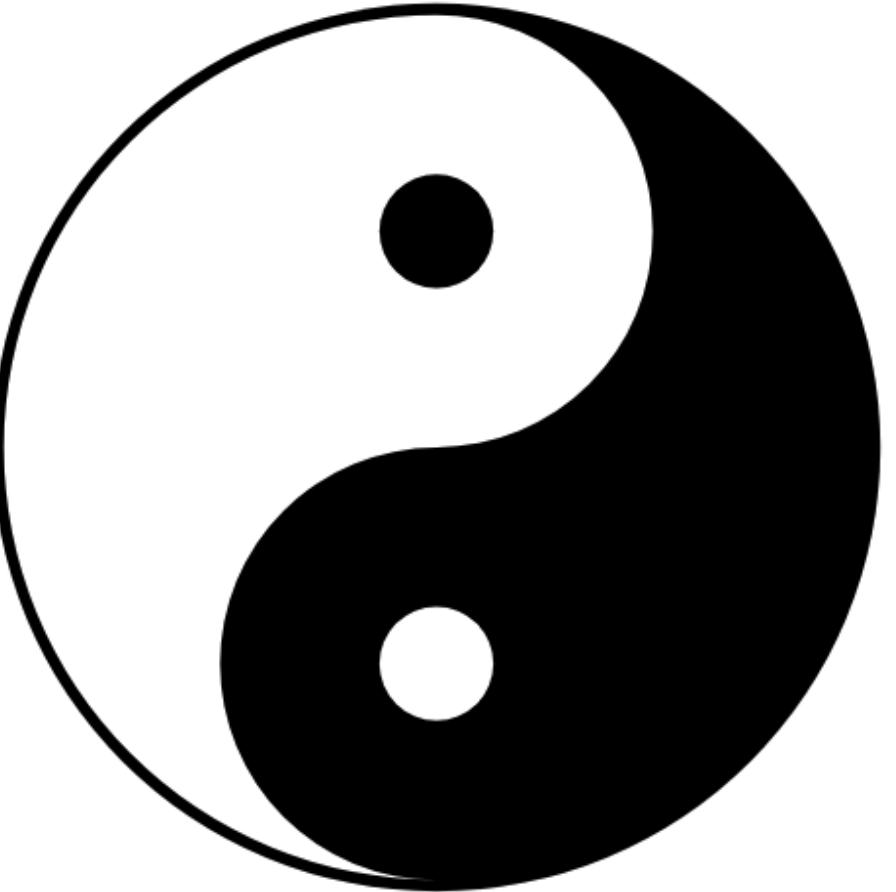
Jeff Bass Bio Bullets (how each ties in to my project)



- Started college pursuing EE, got drafted, fixed radars and meteorological systems in the army (electronics)
- Went to grad school on GI bill; dissertation morphed from econometrics to computer science to a software company (CS)
- Spent 20 years doing stats / data science at big biotech; datasets from genomics to Medicare claims (“Big” data)
- Science: gene sequencing, cell receptor pathway mapping, microbiology, protein chemistry, health care economics (42)
- Small plane pilot; owned a Cessna for a while (“voice ZMQ”)
- Currently retired from income producing endeavors, building the farm, bike touring the California Coast, etc. (lots of time)
- Learning about computer vision, sensors, Raspberry Pi as tools to better manage a permaculture farm (lack of focus)

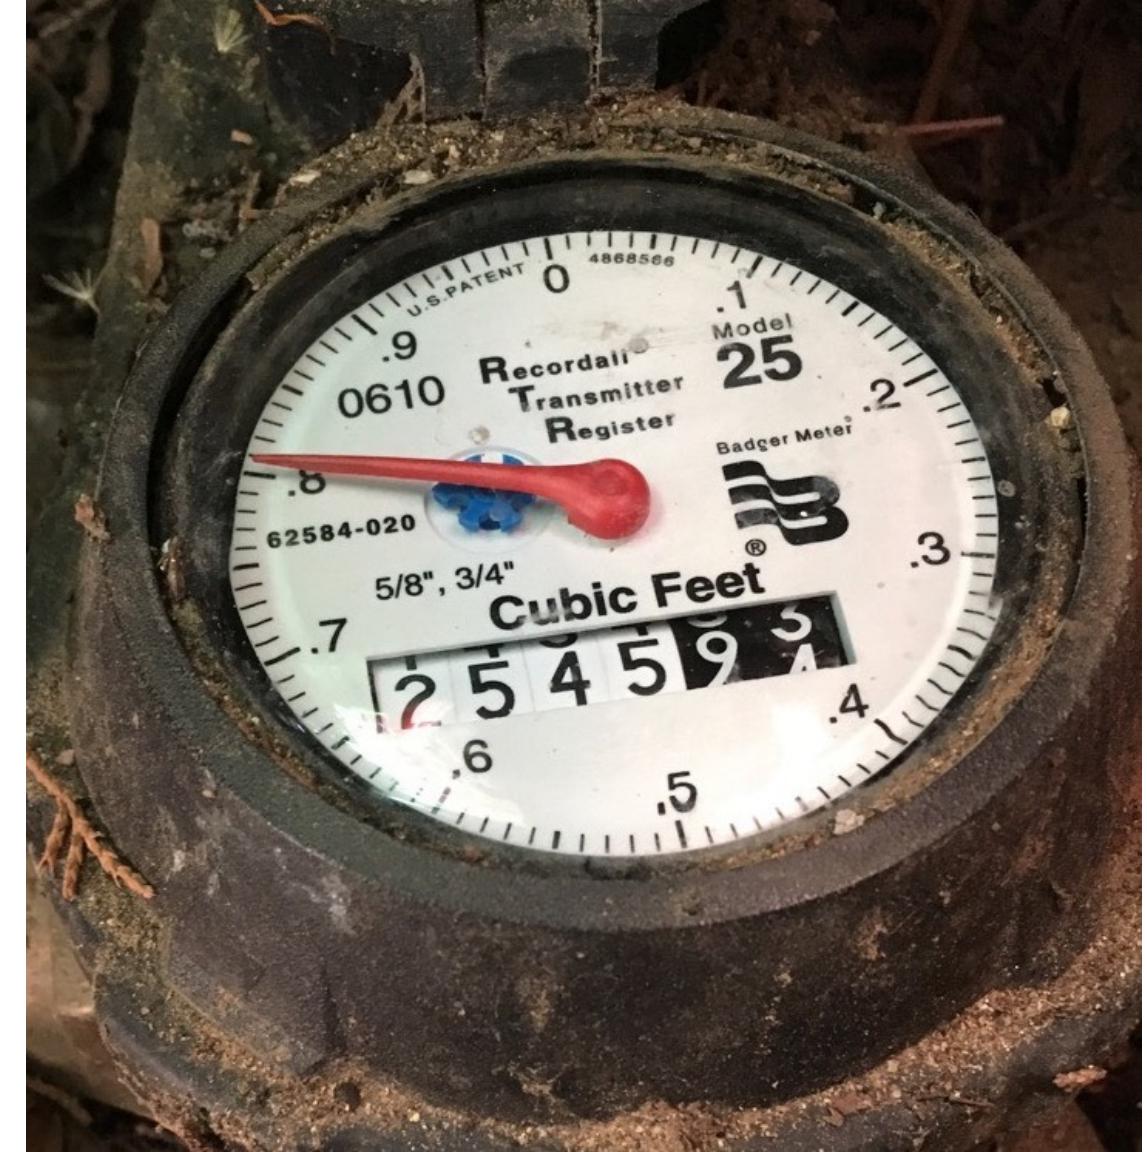
What this talk will cover

- Multi-CPU computer vision pipelines
- Practical applications of CV techniques
- Project evolution and iterative design using Raspberry Pi's (RPi's)
- Yin and Yang of my design decisions; why “this” and “not that”
- But **NOT**:
 - Advanced computer vision or deep learning techniques
 - Deep dive into Permaculture (I can get way too didactic about it)



How Computer Vision helps manage my small permaculture farm

- Reading the water meter; optimizing water use
- Counting bees, butterflies and other pollinators
- Tracking coyotes, rabbits, raccoons and other critters
- Monitoring when garage and barn doors and gates open or close
- Tracking sunlight hours, sunlight intensity, clouds, shadows (photosynthesis strength and availability)
- Monitoring non-camera sensors: temperature, humidity, PIR motion sensors, solar panel power output (large 21KW array)

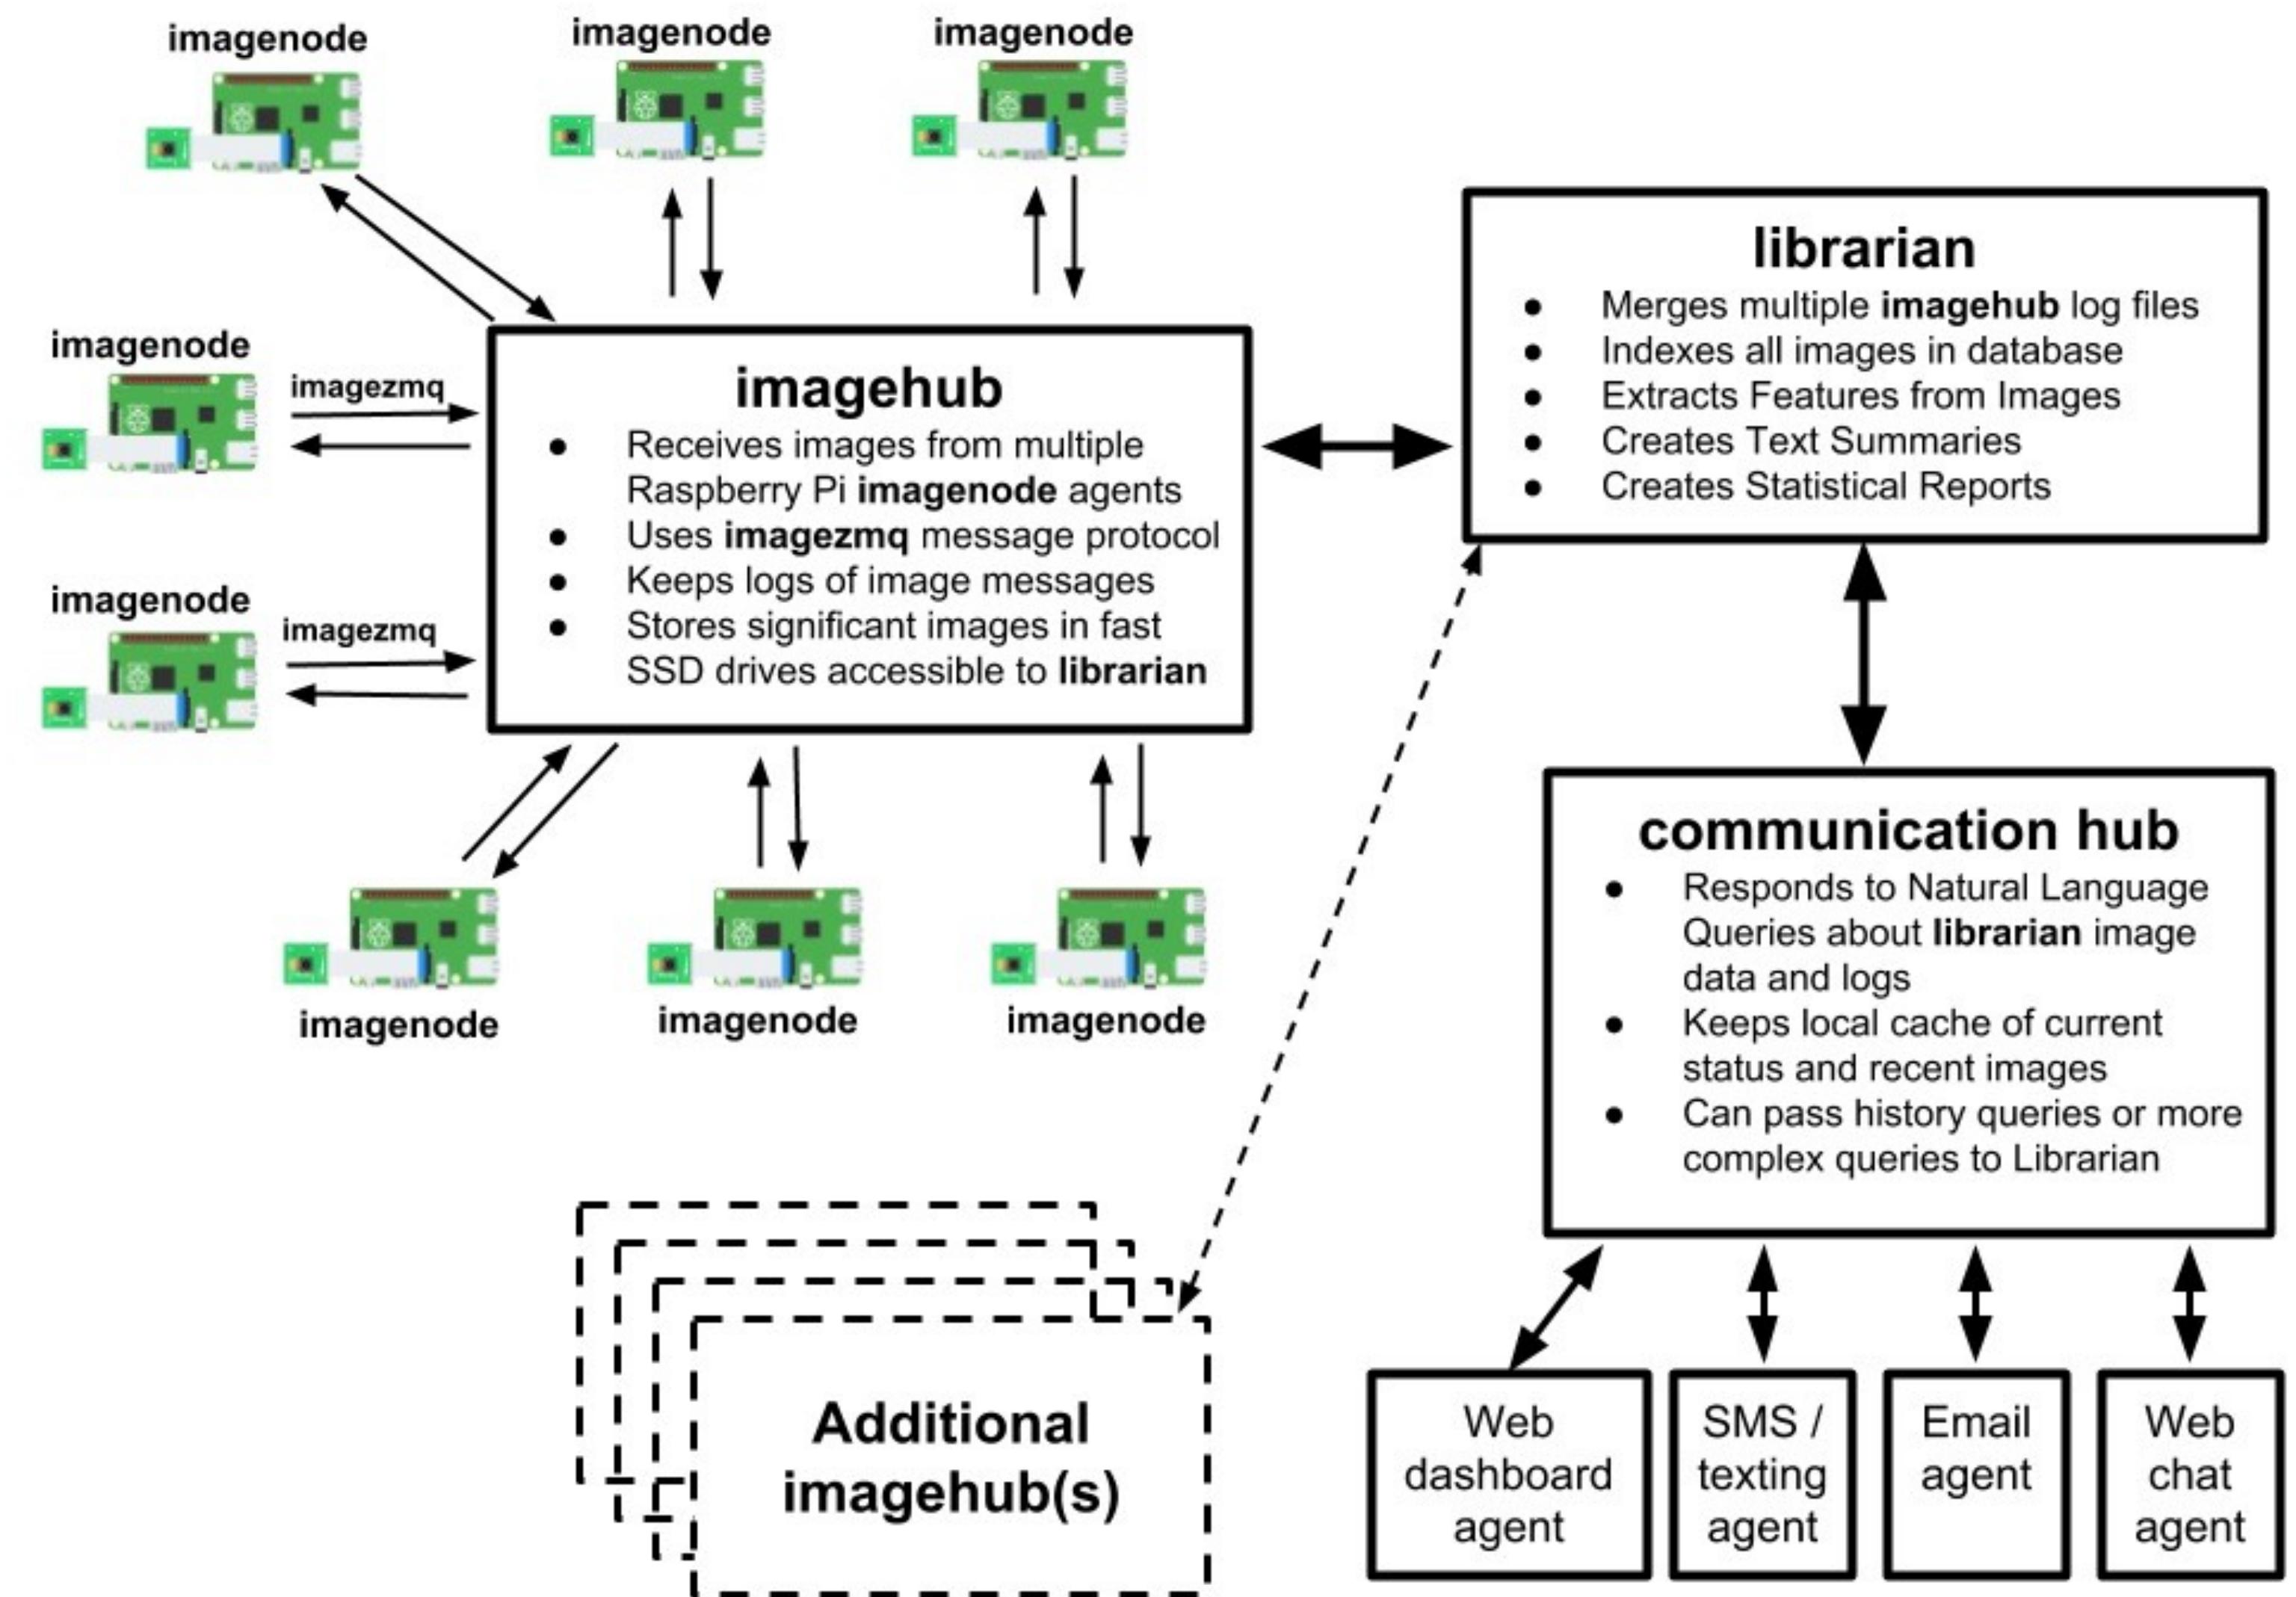


Interlude: What the heck is permaculture?

- Designing for long-term sustainability (centuries rather than decades)
- Creating deep living soil with diverse microbiology (emulating an old growth forest)
- Design choices to prioritize sustainability over efficiency (plant polyculture vs. Big Ag)
- Top down designs for land use, water capture and retention, carbon cycling
- Science based; careful observation; repeatable experiments; open publishing
- Open source seeds, hardware, techniques, designs for disease management
- Farms that are smaller and include native plants, insects, worms and mammals
- Permaculture is “A Revolution Disguised as Organic Gardening”

A Distributed Computer Vision Pipeline

- Raspberry Pi's as nodes
- 1 or 2 cameras per RPi
- 8 to 12 RPi's per Hub
- ZMQ Messaging passes images from RPi's to Hub
- Multiple Hubs feed the Librarian
- Communications Hub



Imagenode <-> imagezmq <-> Imagehub

Pseudocode for imangenode

```
ForeverEventLoop:  
    grab camera image  
    put image in queue  
    analyze queue  
    if SomethingHappened:  
        send event message  
        send images  
        process hub response
```

Pseudocode for imagehub

```
ForeverEventLoop:  
    receive (message, image)  
    store event message  
    store images  
    send response
```

Simple Code – Multiple RPi's to a Single Mac

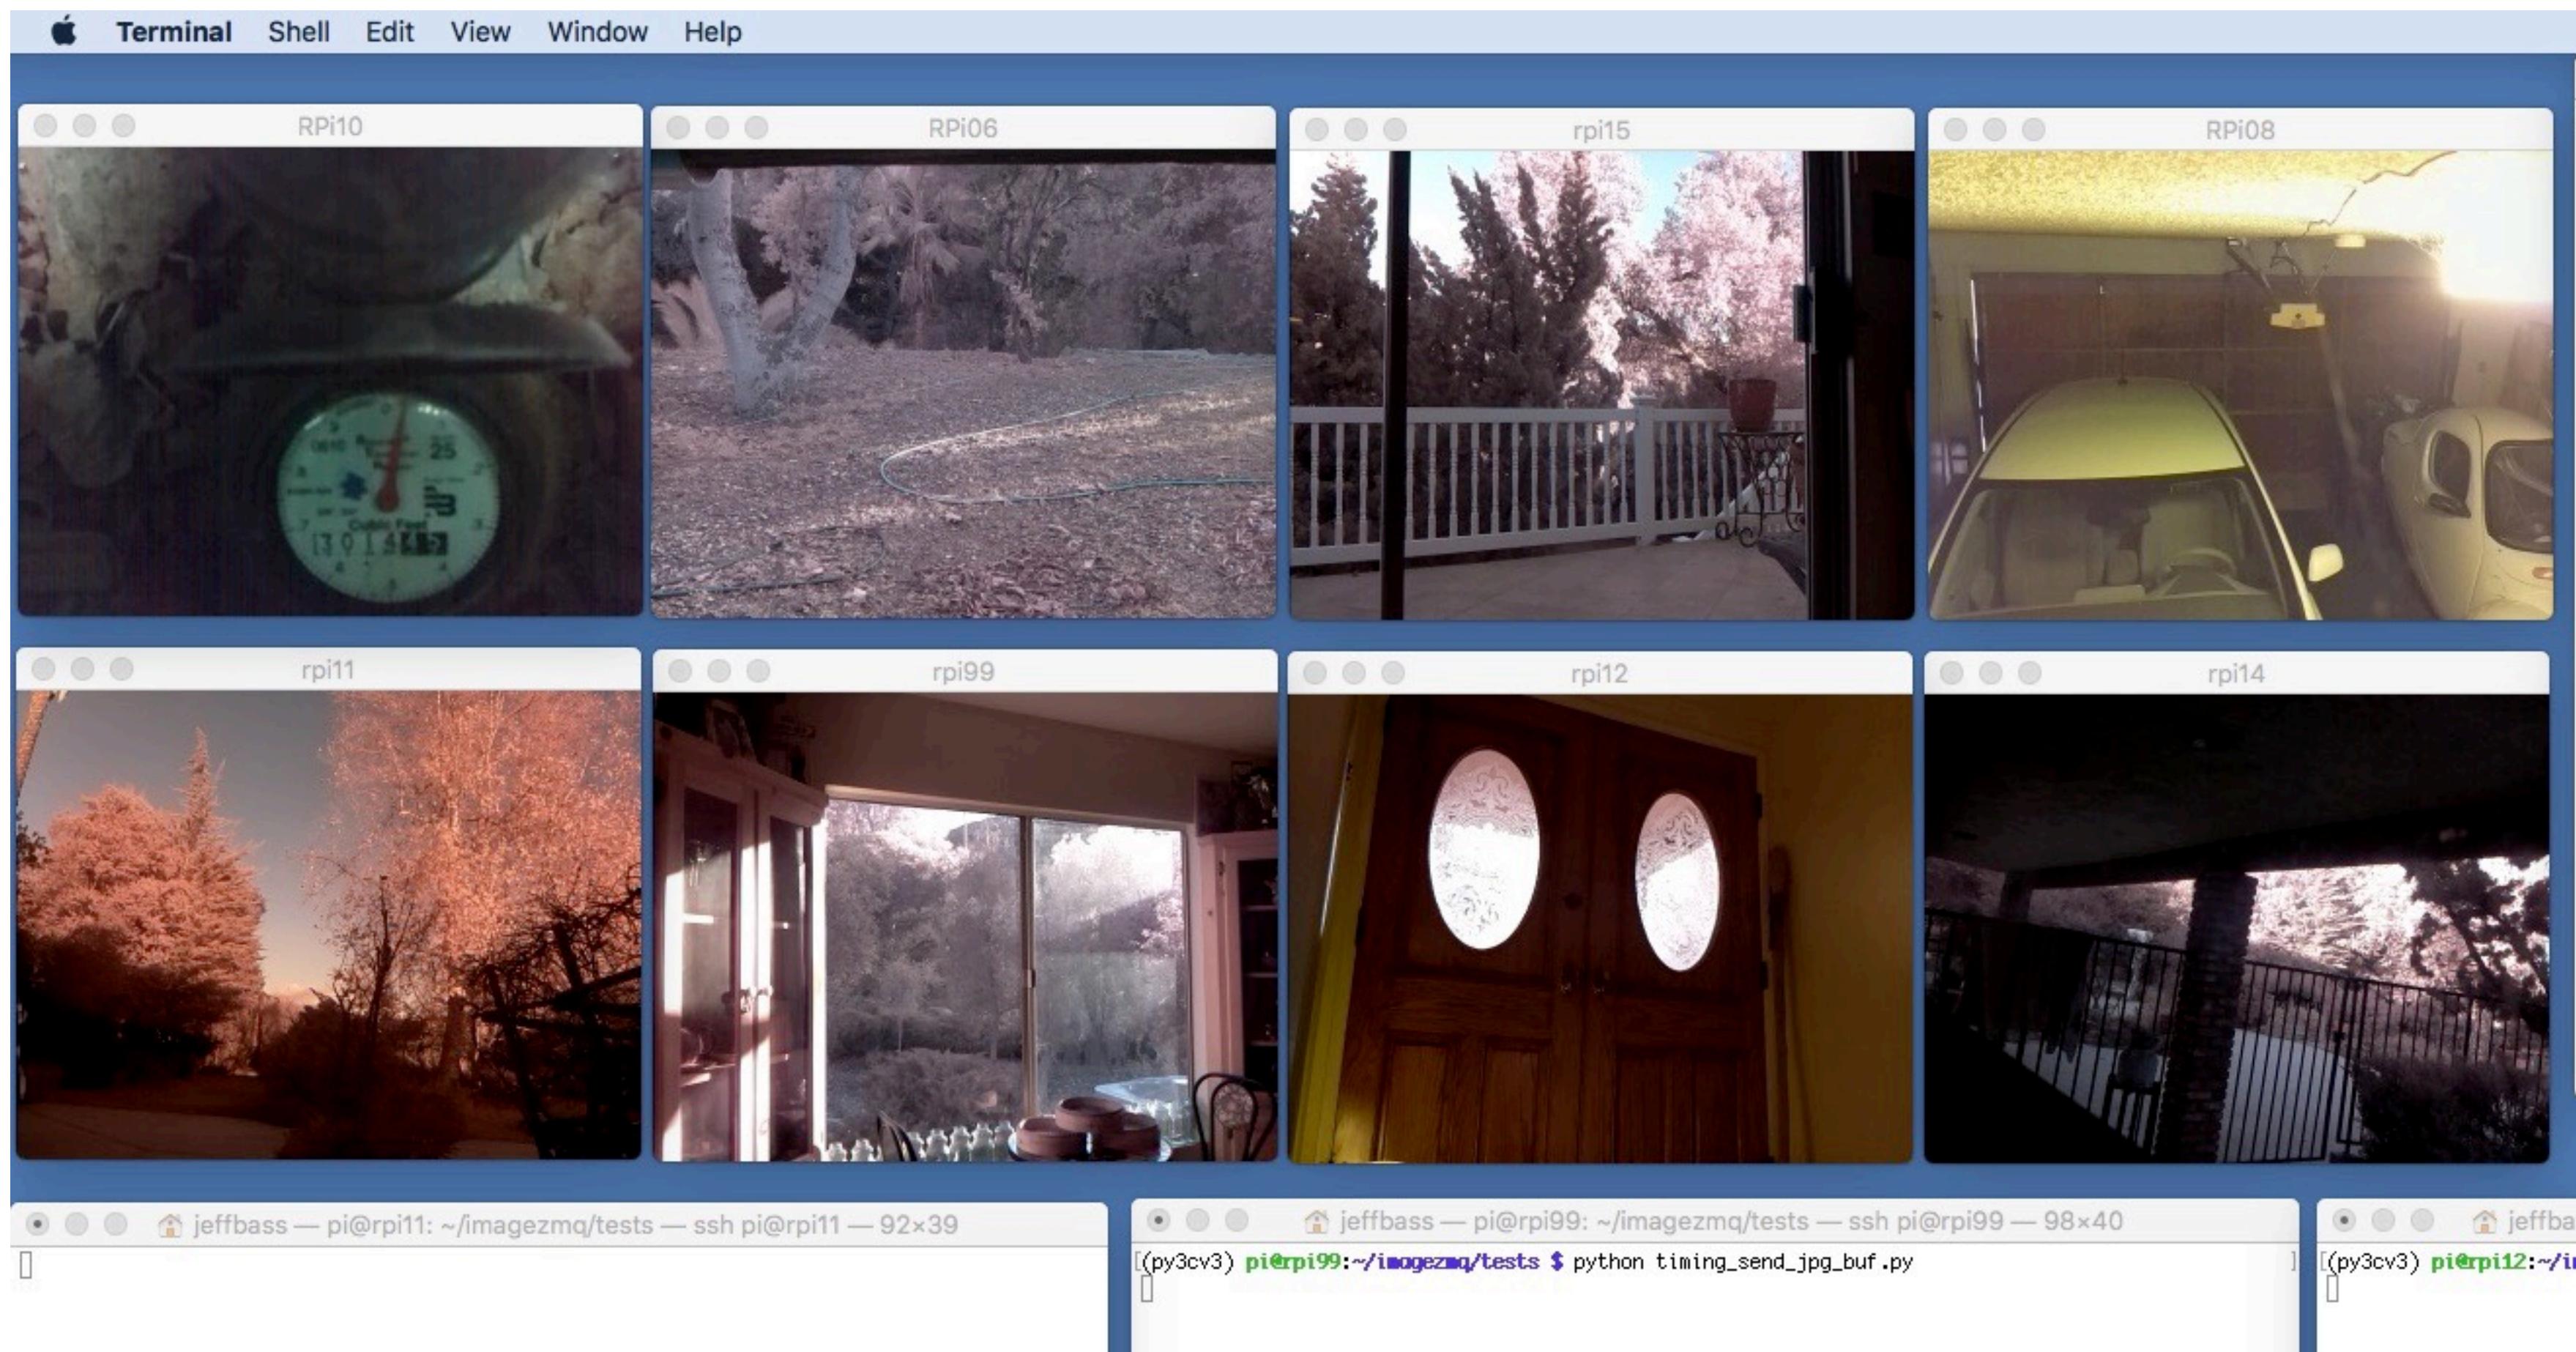
```
1 # run this program on each RPi to send a labelled image stream
2 import socket
3 import time
4 from imutils.video import VideoStream
5 import imagezmq
6
7 sender = imagezmq.ImageSender(connect_to='tcp://jeff-macbook:5555')
8
9 rpi_name = socket.gethostname() # send RPi hostname with each image
10 picam = VideoStream(usePiCamera=True).start()
11 time.sleep(2.0) # allow camera sensor to warm up
12 while True: # send images as stream until Ctrl-C
13     image = picam.read()
14     sender.send_image(rpi_name, image)
```

**Node
Code
(RPi)**

```
1 # run this program on the Mac to display image streams from multiple RPis
2 import cv2
3 import imagezmq
4 image_hub = imagezmq.ImageHub()
5 while True: # show streamed images until Ctrl-C
6     rpi_name, image = image_hub.recv_image()
7     cv2.imshow(rpi_name, image) # 1 window for each RPi
8     cv2.waitKey(1)
9     image_hub.send_reply(b'OK')
```

**Hub
Code
(Mac)**

Simple RPi -> Mac Pipeline: Video Surveillance



An 8 RPi Simulcast in 19 lines of Python (11 per RPi; 8 on a Mac)

Imagenode Code Snippet

```
1 # imagenode.py snippet
2 #      (all the imports, logging setup hidden to fit slide)
3
4 settings = Settings() # get settings for node cameras, ROIs, GPIO
5 node = ImageNode(settings) # start ZMQ, cameras and other sensors
6
7 # forever event loop
8 while True:
9     node.get_sensor_data() # grab camera images and other sensor data
10    node.process_sensor_data() # detect motion, etc.
11    while len(node.msg_q) > 0:
12        try:
13            with Patience(settings.patience): # wait for hub response
14                text, image = node.msg_q.popleft()
15                hub_reply = node.send_frame(text, image)
16            except Patience.Timeout: # except no timely response from hub
17                hub_reply = node.fix_comm_link()
18            node.process_hub_reply(hub_reply)
19 # .....
```

Example: Reading a Water Meter

What the RPi does:

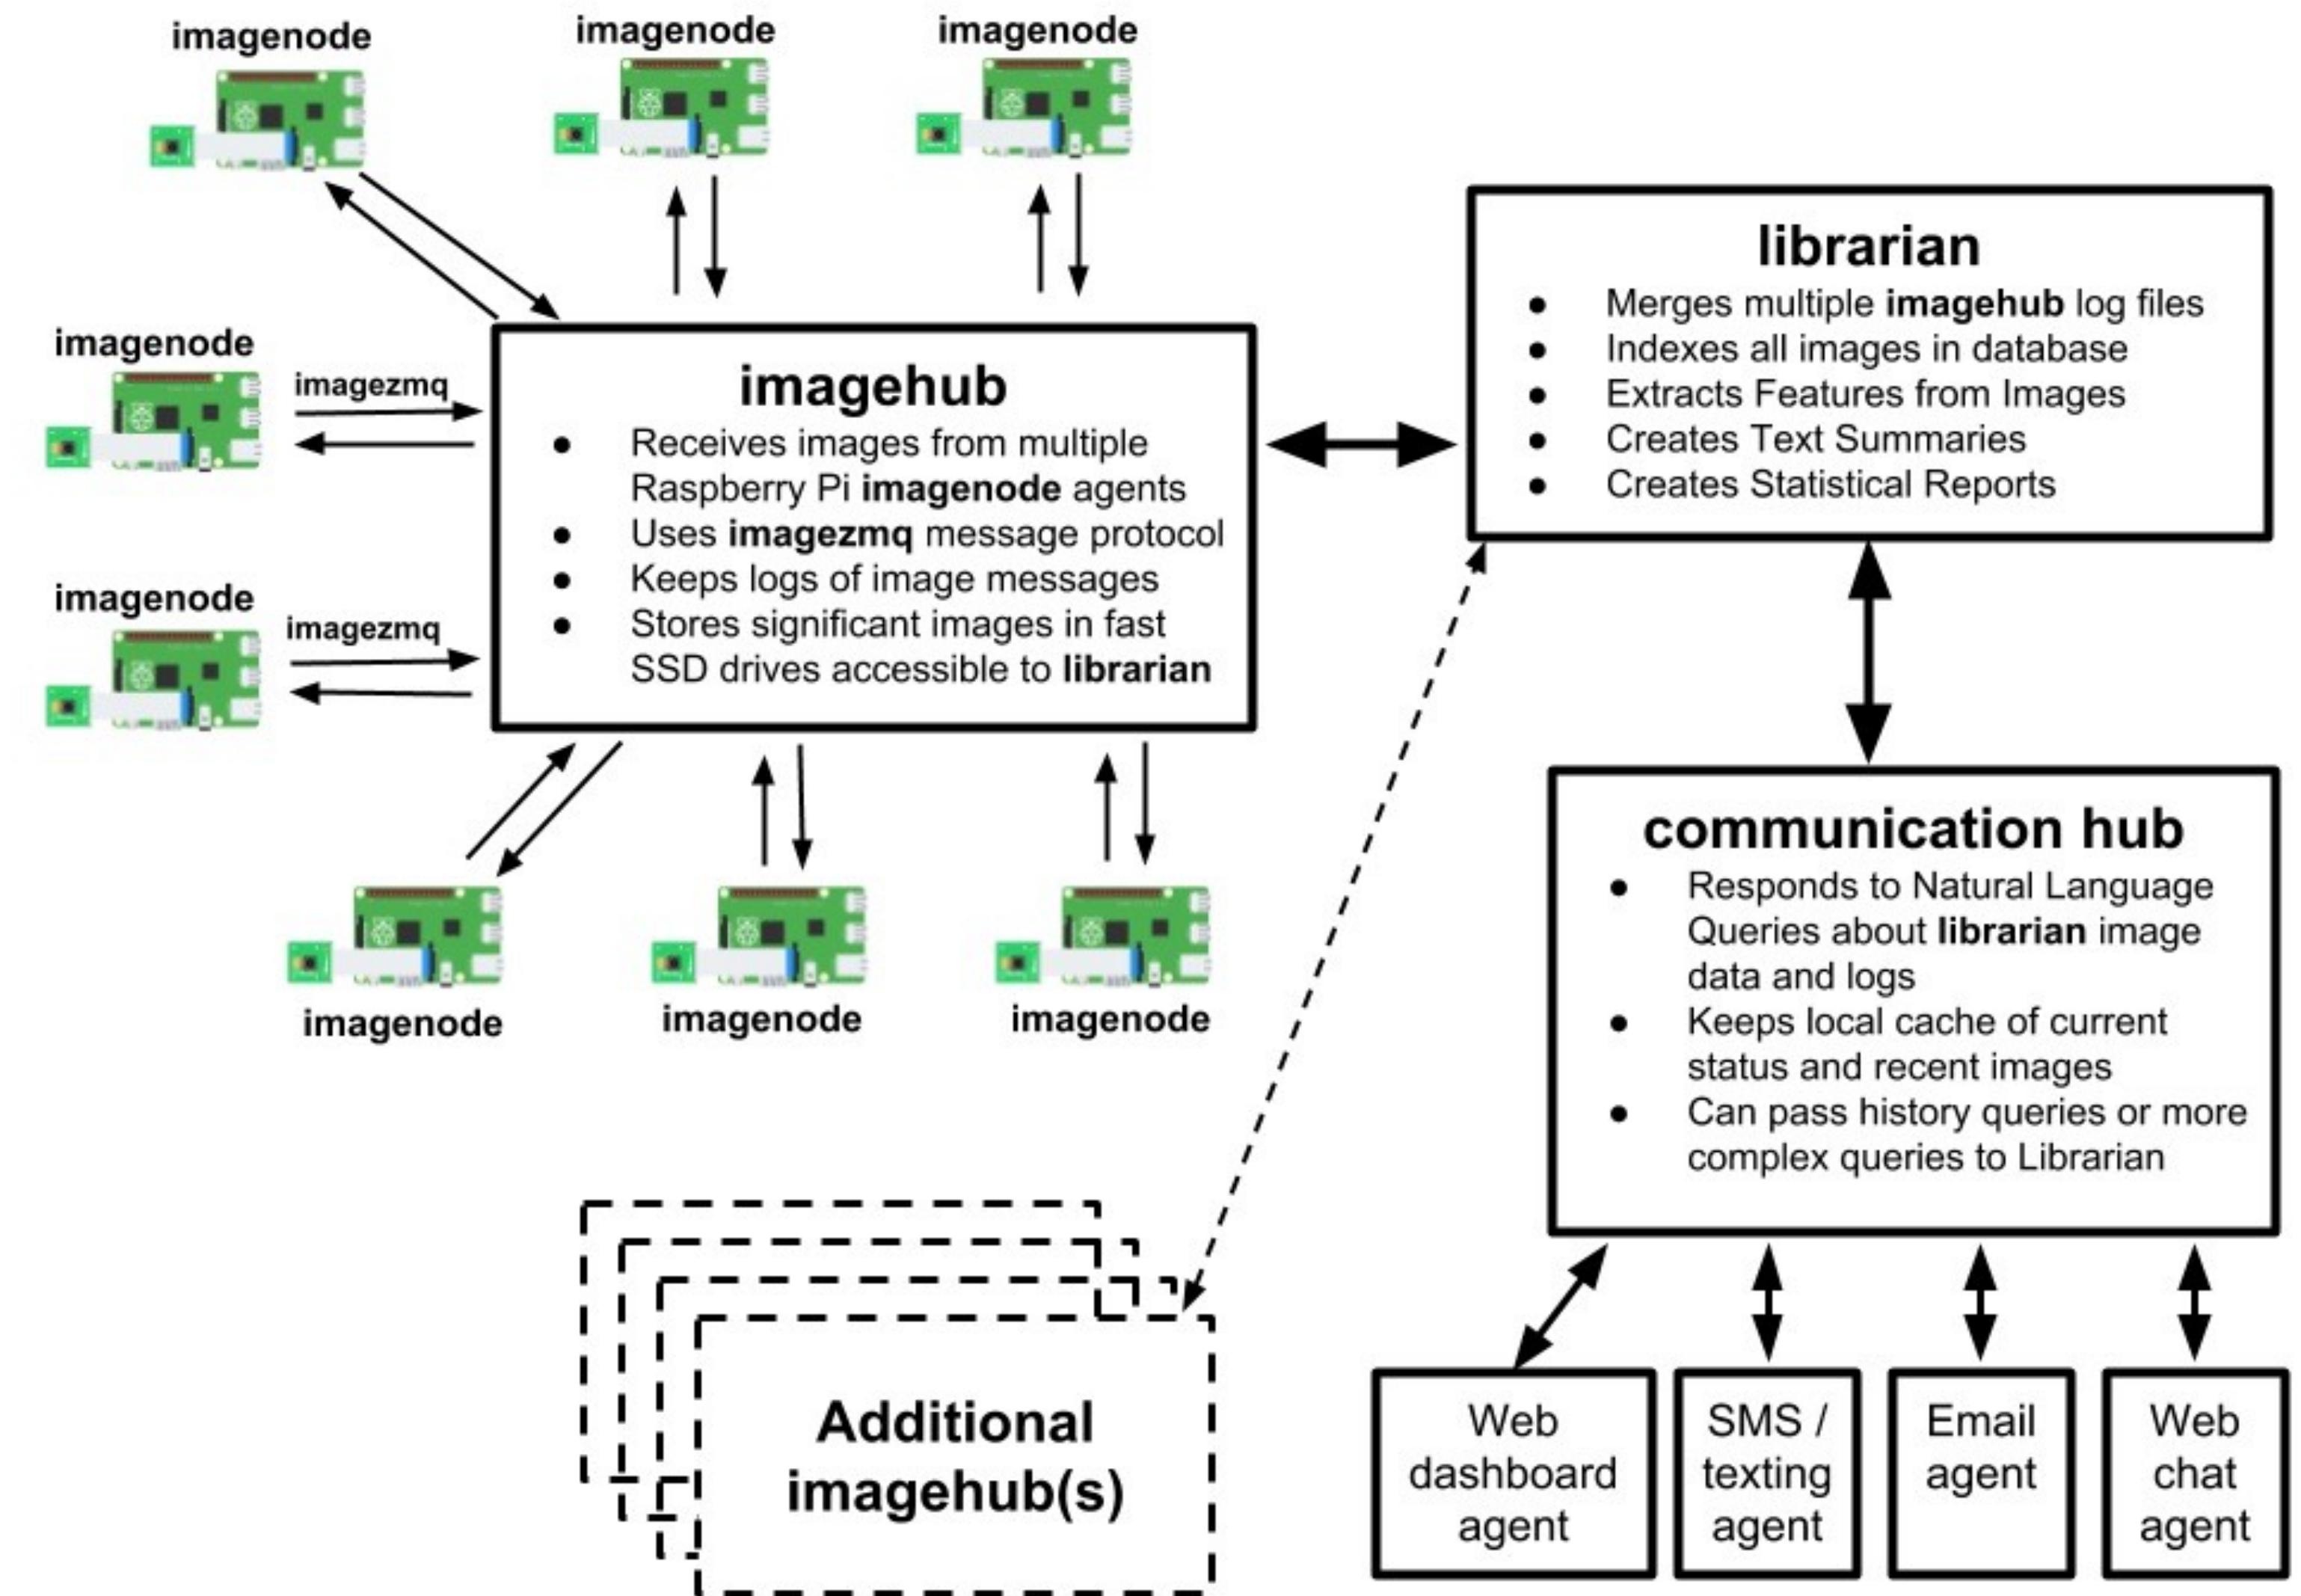
- Turns on LEDs to light the meter
- Continuously grabs frames
- Does ROI cropping, conversion to grayscale, thresholding, dilation
- Detects needle & leak spinner motion
 - Sends event msg (“water flowing”)
 - Sends images of event
- Xfers images & messages via ZMQ

What the Mac does:

- Acts as ZMQ hub for multiple RPi's
- Receives and acknowledges images and messages from multiple RPi's
- Adds events messages to events log
- Stores and indexes images
- Extracts digits to read the meter
- Keeps history of events; does stats
- Responds to queries about status

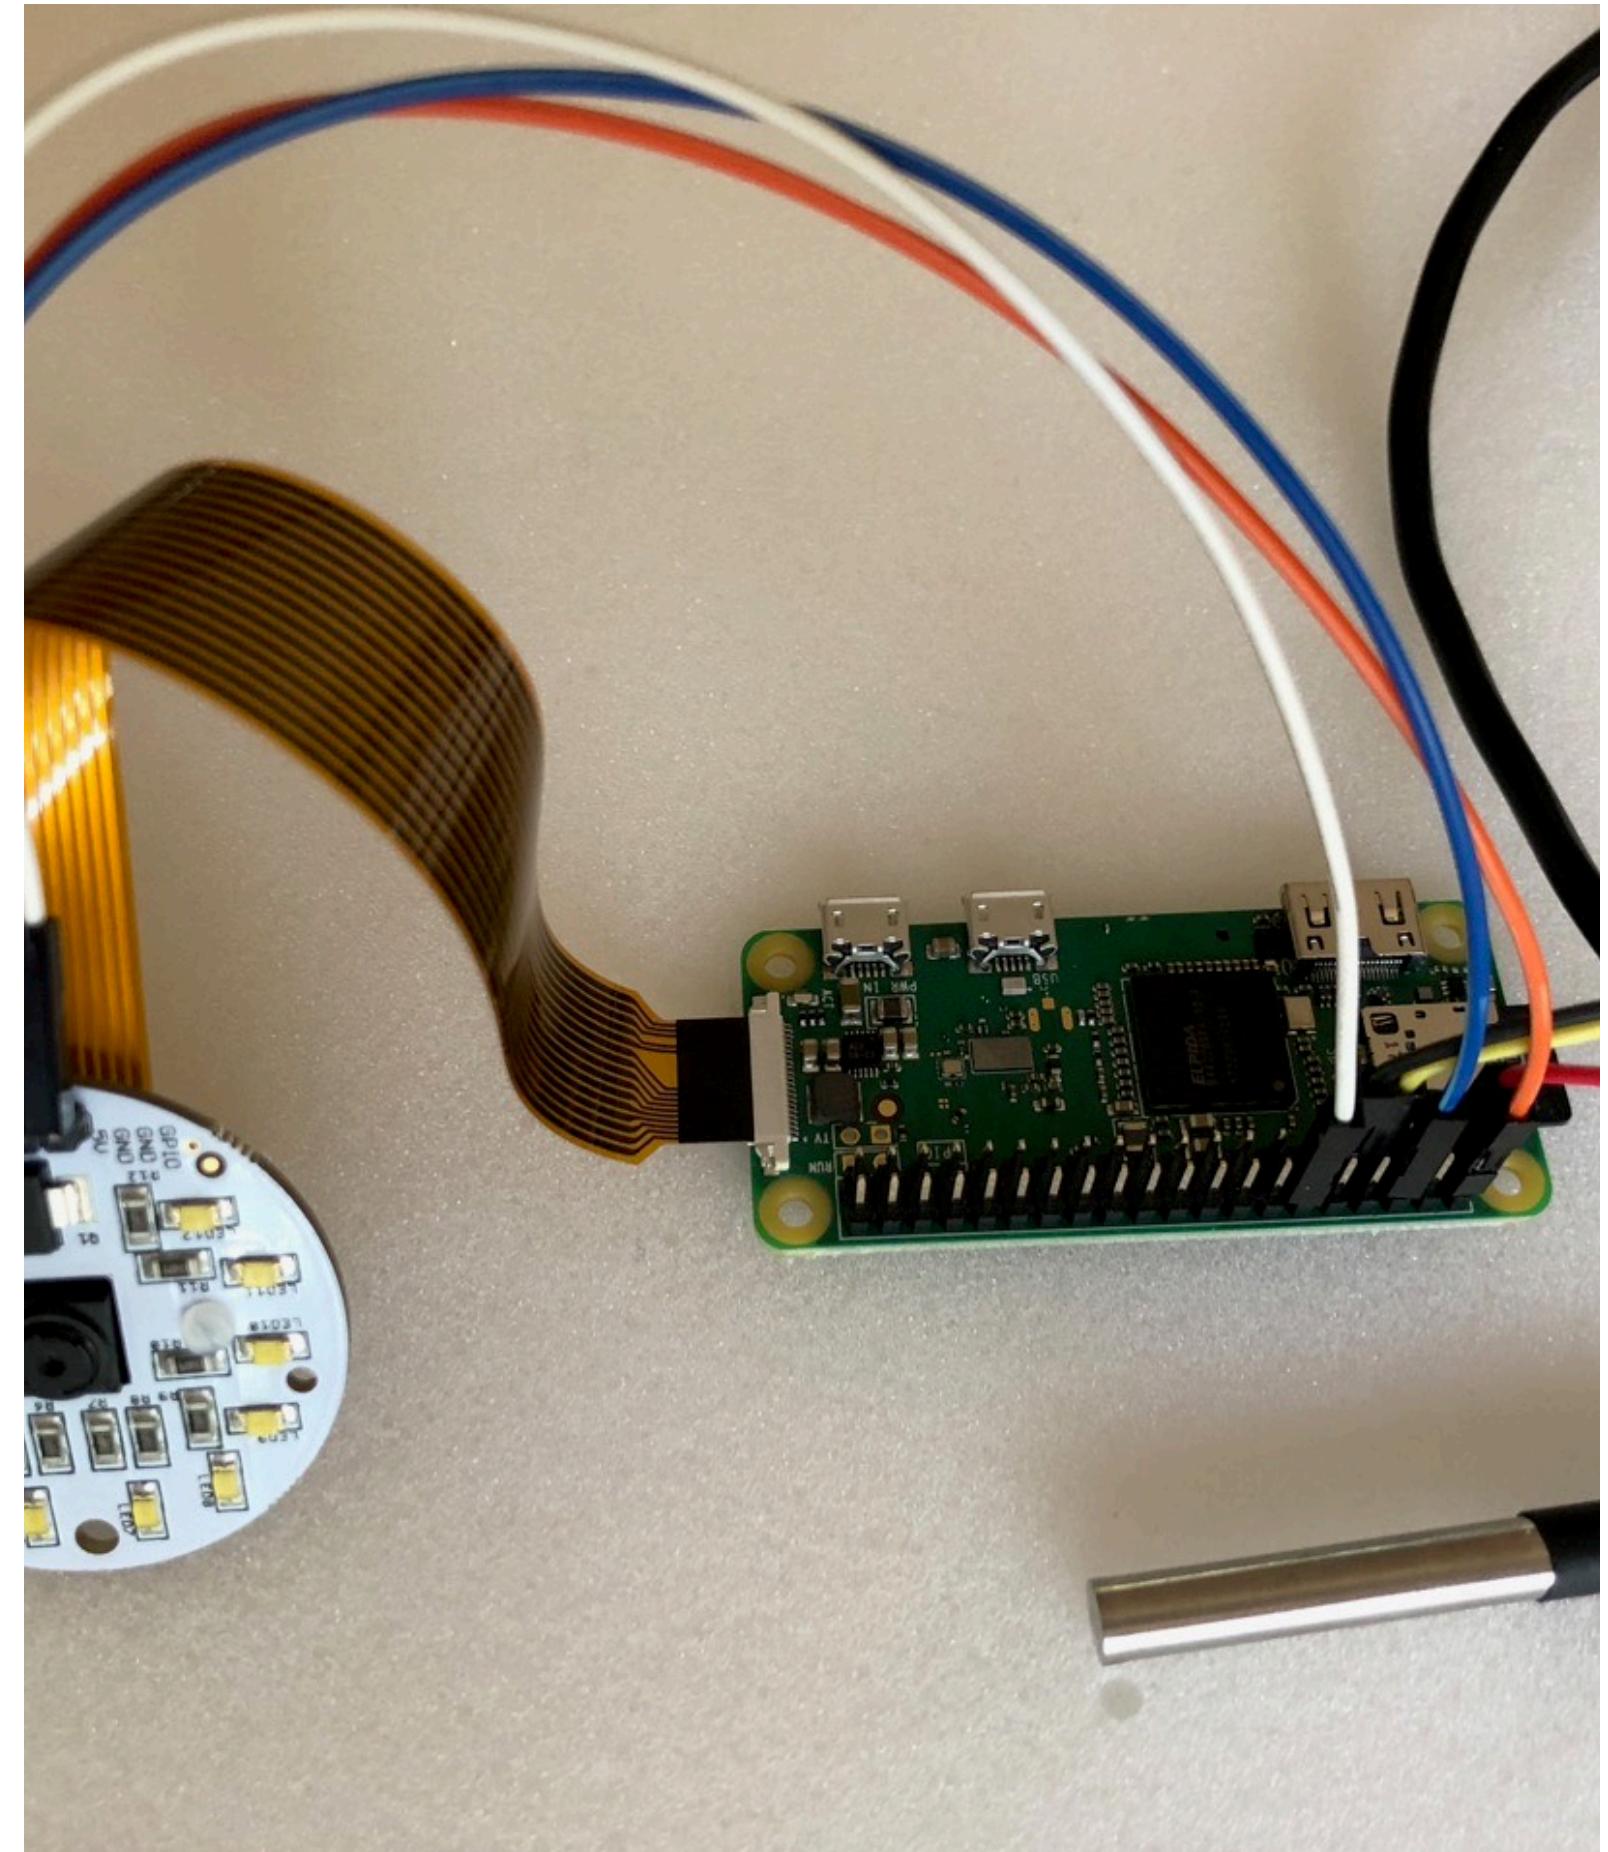
Distributed Computer Vision Pipeline (revisited)

- Raspberry Pi's as nodes
- 1 or 2 cameras per RPi
- 8 to 12 RPi's per Hub
- ZMQ Messaging passes images from RPi's to Hub
- Multiple Hubs feed the Librarian
- Communications Hub

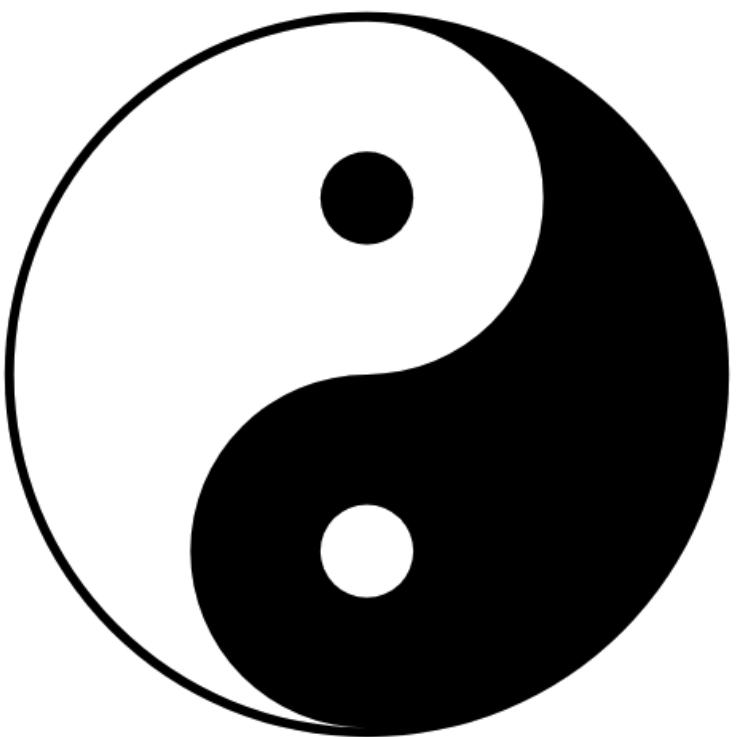


Raspberry Pi's are Great (at some things)

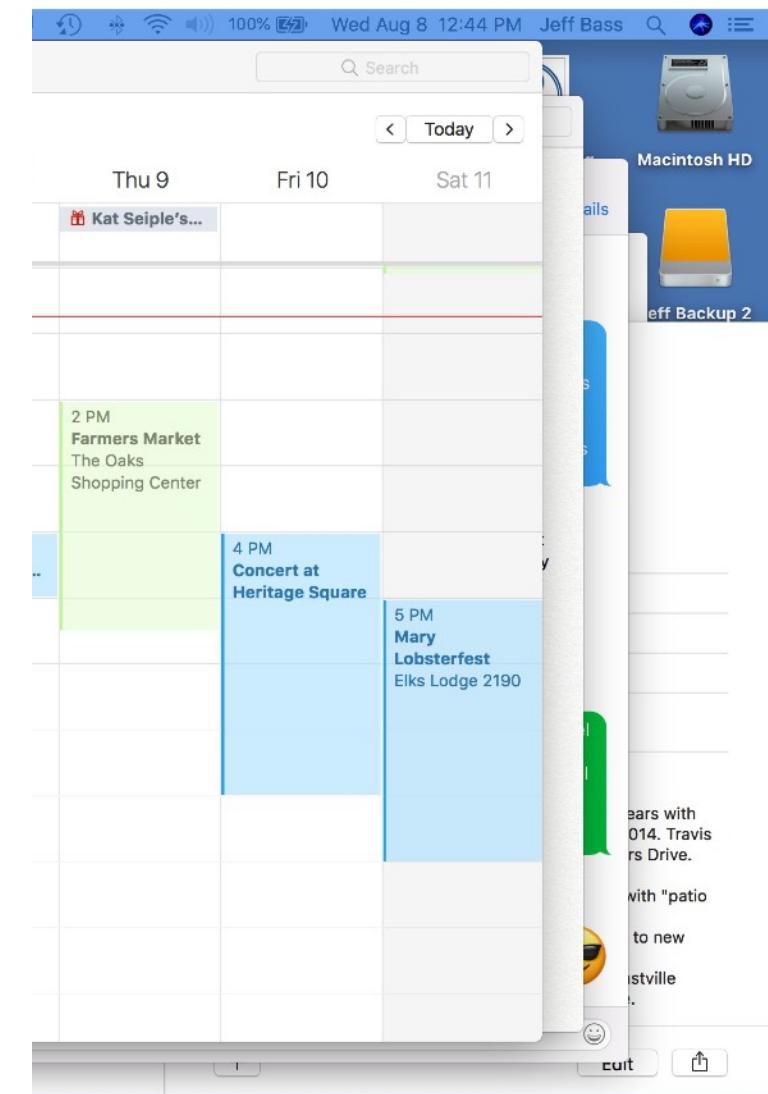
- Cheap, reliable, **fanless**, dust tolerant
- Physically small (Pi Zero) for sensor enclosures
- Very low power; quick restart after power failures
- PiCamera has helpful adjustable settings
- **BUT...**
 - SD Cards not reliable for writing large image files
 - SD Card writing is **VERY** slow
 - Limited memory and limited processing capacity
 - Slower USB bus: slow Ethernet; slow WiFi
 - USB webcams on RPi are slow and cranky



Macs are Great (at some things)



- Fast WiFi & Ethernet
- Fast SSD storage
(for lots of images)
- More computing power for elaborate processing of images
- Digits recognition; feature extraction; classification
- **But...**one Mac laptop uses as much power as 8 or 10 RPis
- Macs are expensive and not suitable for outdoors

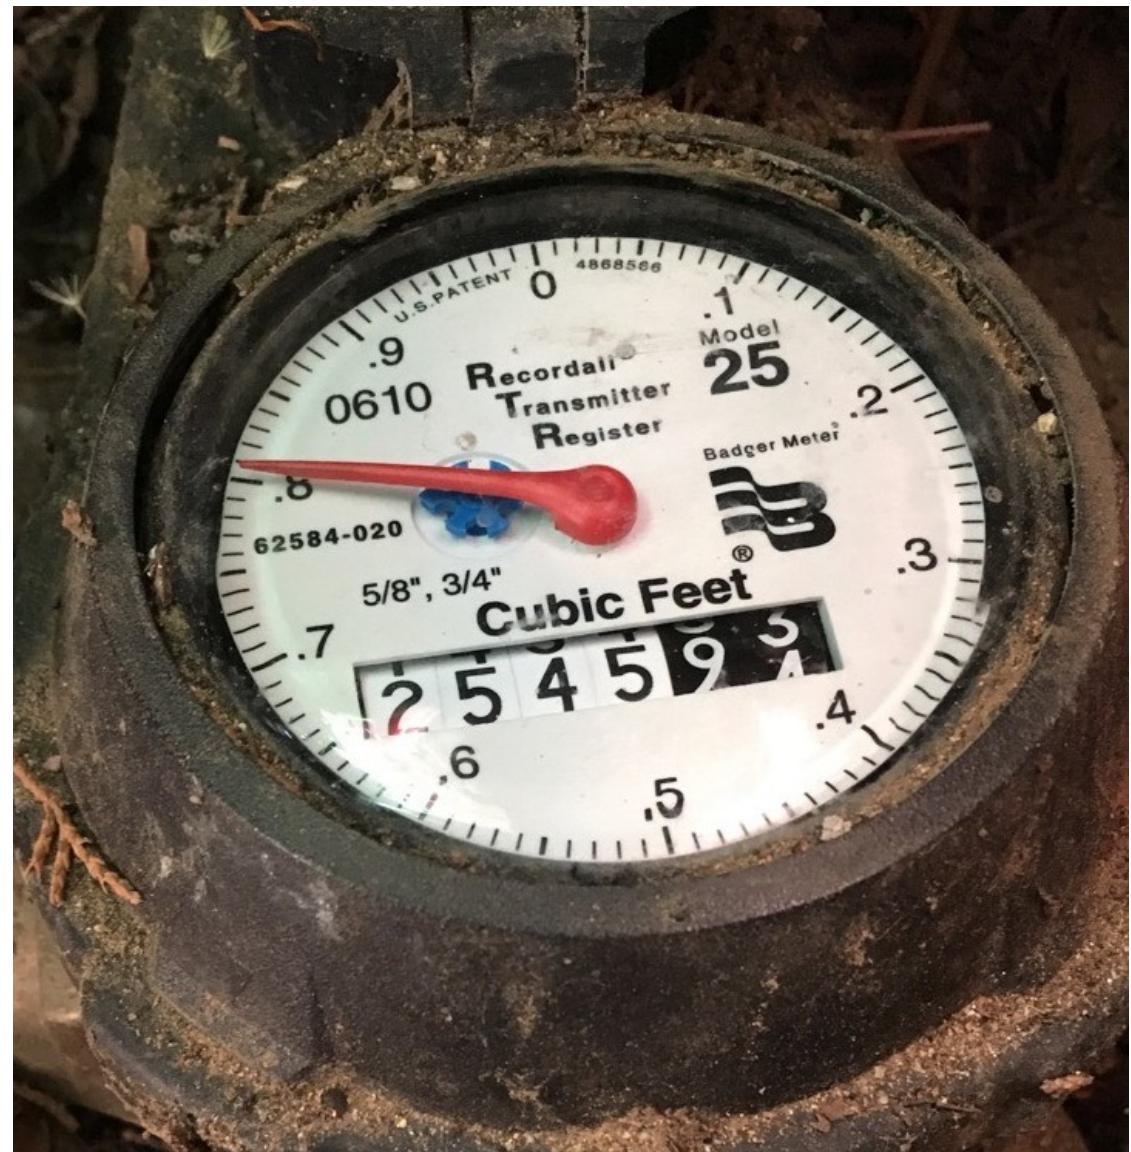


```
Terminal Shell Edit View Window Help
Project
  imagenode.py imaging.py imangenode.yaml settings-yaml.rst nodehealth.rst testing.rst README.rst imangenode-uses.rst utils.py nodehealth.py Code Bits.py
  tests -- jeffbass@jeff-macbook: ~/SDBdev2/imagenode
  # Settings yaml file for imangenode.py PiCamera
  node:
    name: WaterMeter
    queuebox: 50
    patience: 10
    hardware: 10
    send_type: jpg
    print_settings: False
  hub_address:
    H1: tcp://jeff-macbook:5555
    H2: tcp://192.168.1.55:5555
  cameras:
    P1:
      viewname: MotorSkill
      resolution: (640, 480)
      framerate: 16
      vflip: True
      send_count: 1
      event: 50
      detectors:
        motion: (flowing, off)
      ROI1:
        width: 95
        top: 45
        bottom: 90
        left: 40
        right: 82
        delta_thresh: 5
        spread: 1
        min_area: 35
      lights:
        L1:
          home: LEDArray
          gpio: 18
      sensors:
        T1:
          type: DS18B20
          gpio: 4
    ~
  "watermeter.yaml" 39L, 683C written

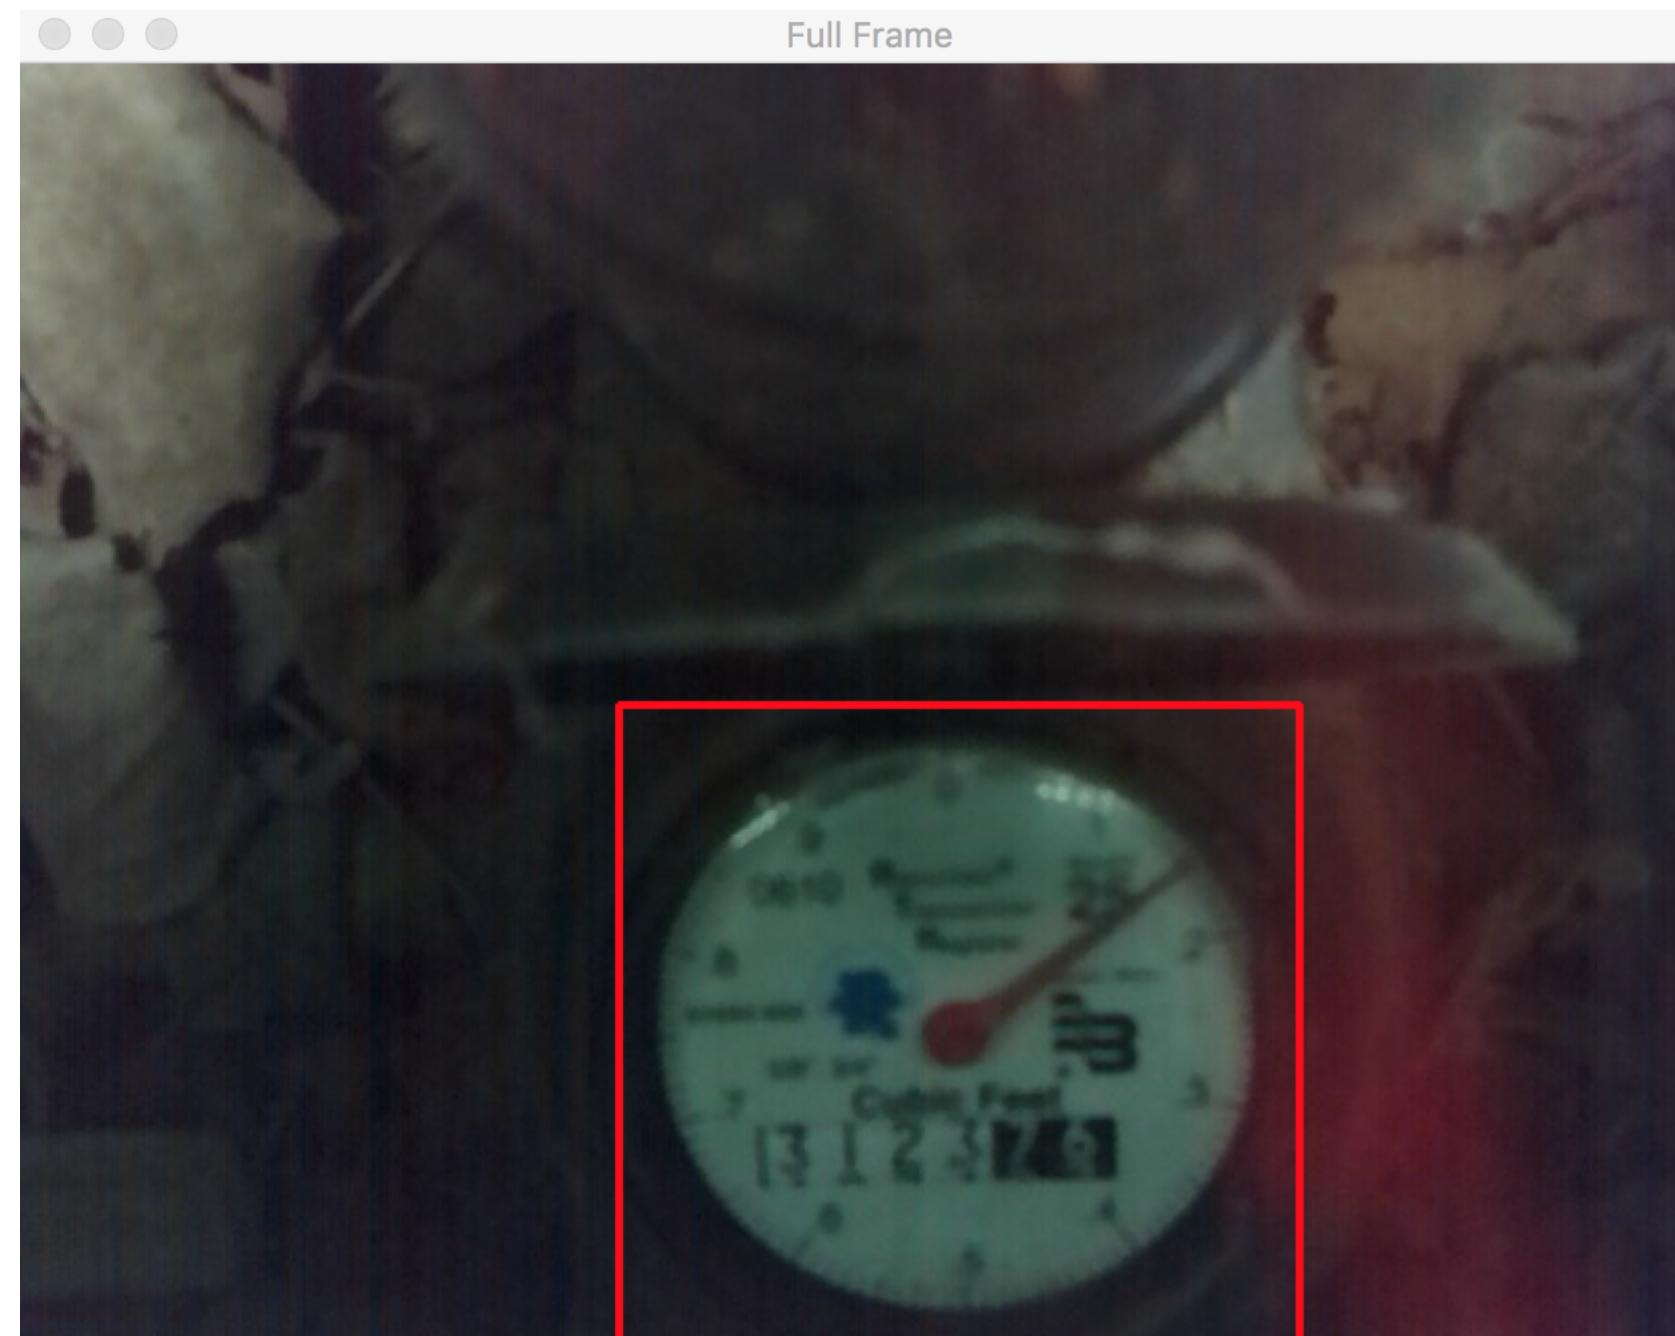
16 import pprint
17 import logging
18 import logging.handlers
19 import traceback
20 from tools.utils import clean_shutdown_when_killed
21 from tools.utils import Patience
22 from tools.imaging import Settings
23 from tools.imaging import ImageNode
24
25 def main():
26     # set up controlled shutdown when Kill Process or SIGTERM received
27     signal.signal(signal.SIGTERM, clean_shutdown_when_killed)
28
29     # start logging
30     log = logging.getLogger()
31     handler = logging.handlers.RotatingFileHandler('imagenode.log',
32         maxBytes=15000, backupCount=5)
33     formatter = logging.Formatter('%(asctime)s ~ %(message)s')
34     handler.setFormatter(formatter)
35     log.addHandler(handler)
36     log.setLevel(logging.DEBUG)
37
38     try:
39         log.info('Starting imangenode.py')
40         settings = Settings() # get settings for node cameras, ROIs,
41         print('Settings node name:', settings.nodename)
42         node = ImageNode(settings) # start ZMQ, cameras and other sensors
43         print('System Type is', node.health.sys_type)
44         # forever event loop
45         while True:
46             node.get_sensor_data() # grab camera images and other sensor data
47             node.process_sensor_data() # detect motion, etc.
48             while len(node.msg_q) > 0:
49                 try:
50                     with Patience(settings.patience): # wait for hub response
51                         # hub_reply = node.send_sensor_data()
52                         text, image = node.msg_q.popleft()
53                         hub_reply = node.send_frame(text, image)
54                 except Patience.Timeout: # except no timely response from hub
55                     log.info("No sensor hub reply for "
56                             + str(int(settings.patience)) + ' seconds')
57                     hub_reply = node.fix_comm_link()
58                     node.process_hub_reply(hub_reply)
59
60             except (KeyboardInterrupt, SystemExit):
61                 log.warning('Ctrl-C was pressed or SIGTERM was received.')
62             except Exception as ex: # traceback will appear in log
63                 log.exception('Unanticipated error with no Exception handler.')
64                 print('Traceback error:', ex)
65             finally:
66                 if 'node' in locals():
67                     node.closeall() # close cameras, GPIO0, files
68                     log.info('Exiting imangenode.py')
69                     sys.exit()
70
71     if __name__ == '__main__':
72         main()

Imagenode/imangenode.py 46:51
```

Image Pipeline for Reading a Water Meter

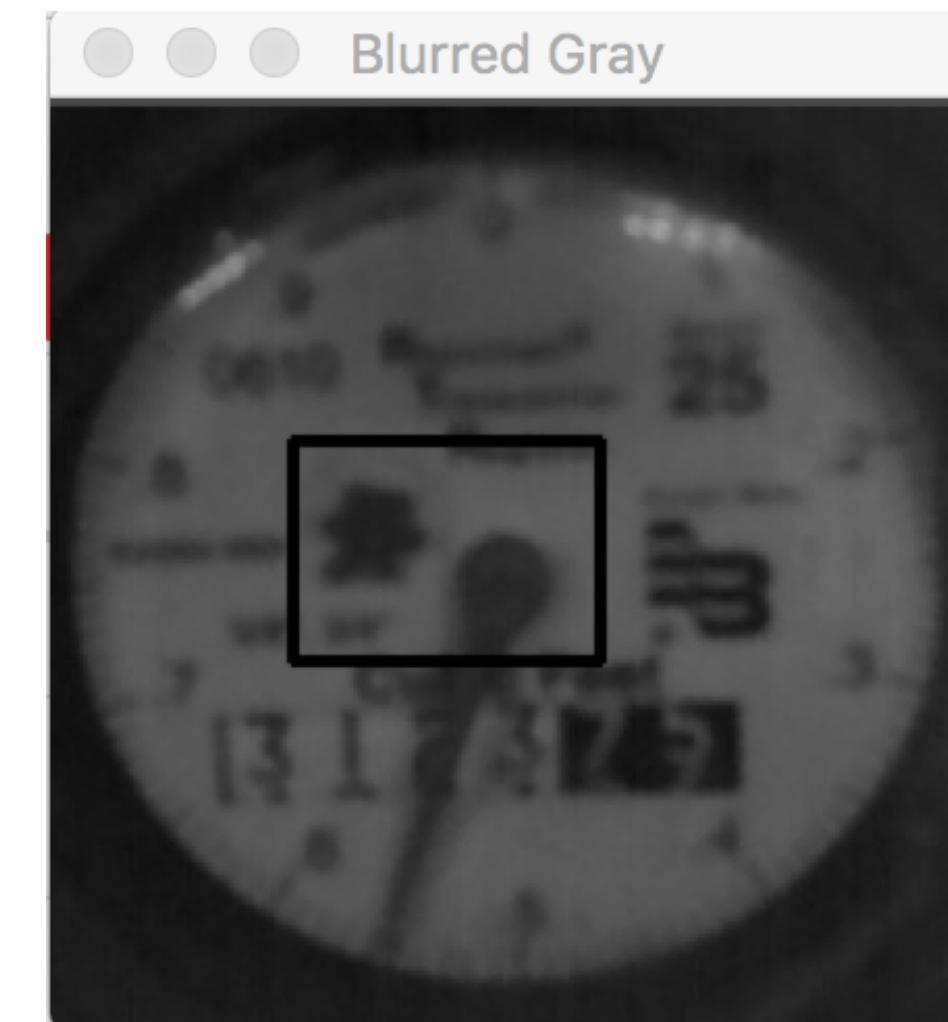


**Daylight view
with the water
meter cover off**



**Full Frame
from RPi PiCamera.
Lit with 24 low power
LEDs to minimize
reflections and glare**

(Gotcha: rotating needle can occlude digits)



**RPi does
frame to frame
differencing
to detect motion
of spinning bits
in the rectangle**



**Mac does
digit masking,
digit classification
to read the
value of the
digits**

There are many, many CV settings...

Need to iteratively optimize all settings

- Settings for:
 - Node name, queue sizes, wait times
 - Hub addresses for multiple hubs
 - Cameras
 - Viewname, resolution, framerate, send quantity
 - Detectors (motion, light, color, etc)
 - ROI dimensions and corners
 - Detection parameters (delta_threshold, area)
 - Lights (under GPIO control)
 - Sensors (temperature, humidity, PIR)
- Using YAML files for settings
 - Feels more Pythonic than json or configs
 - Large settings file becomes a nested Python dict

```
# Settings yaml file for imangenode.py PiCamera
---
node:
  name: WaterMeter
  queue_max: 50
  patience: 10
  heartbeat: 10
  send_type: jpg
  print_settings: False
hub_address:
  H1: tcp://jeff-macbook:5555
  H2: tcp://192.168.1.155:5555
cameras:
  P1:
    viewname: MeterDial
    resolution: (640, 480)
    framerate: 16
    vflip: True
    send_amount:
      event: 50
    detectors:
      motion: (flowing, off)
      ROI1:
        width: 95
        top: 45
        bottom: 90
        left: 40
        right: 82
        delta_thresh: 5
        spread: 1
        min_area: 35
    lights:
      L1:
        name: LEDarray
        gpio: 18
    sensors:
      T1:
        type: DS18B20
        gpio: 4
      ~
      ~
"watermeter.yaml" 39L, 683C written
```

39,1

All

Communications: ZMQ



Distributed Messaging

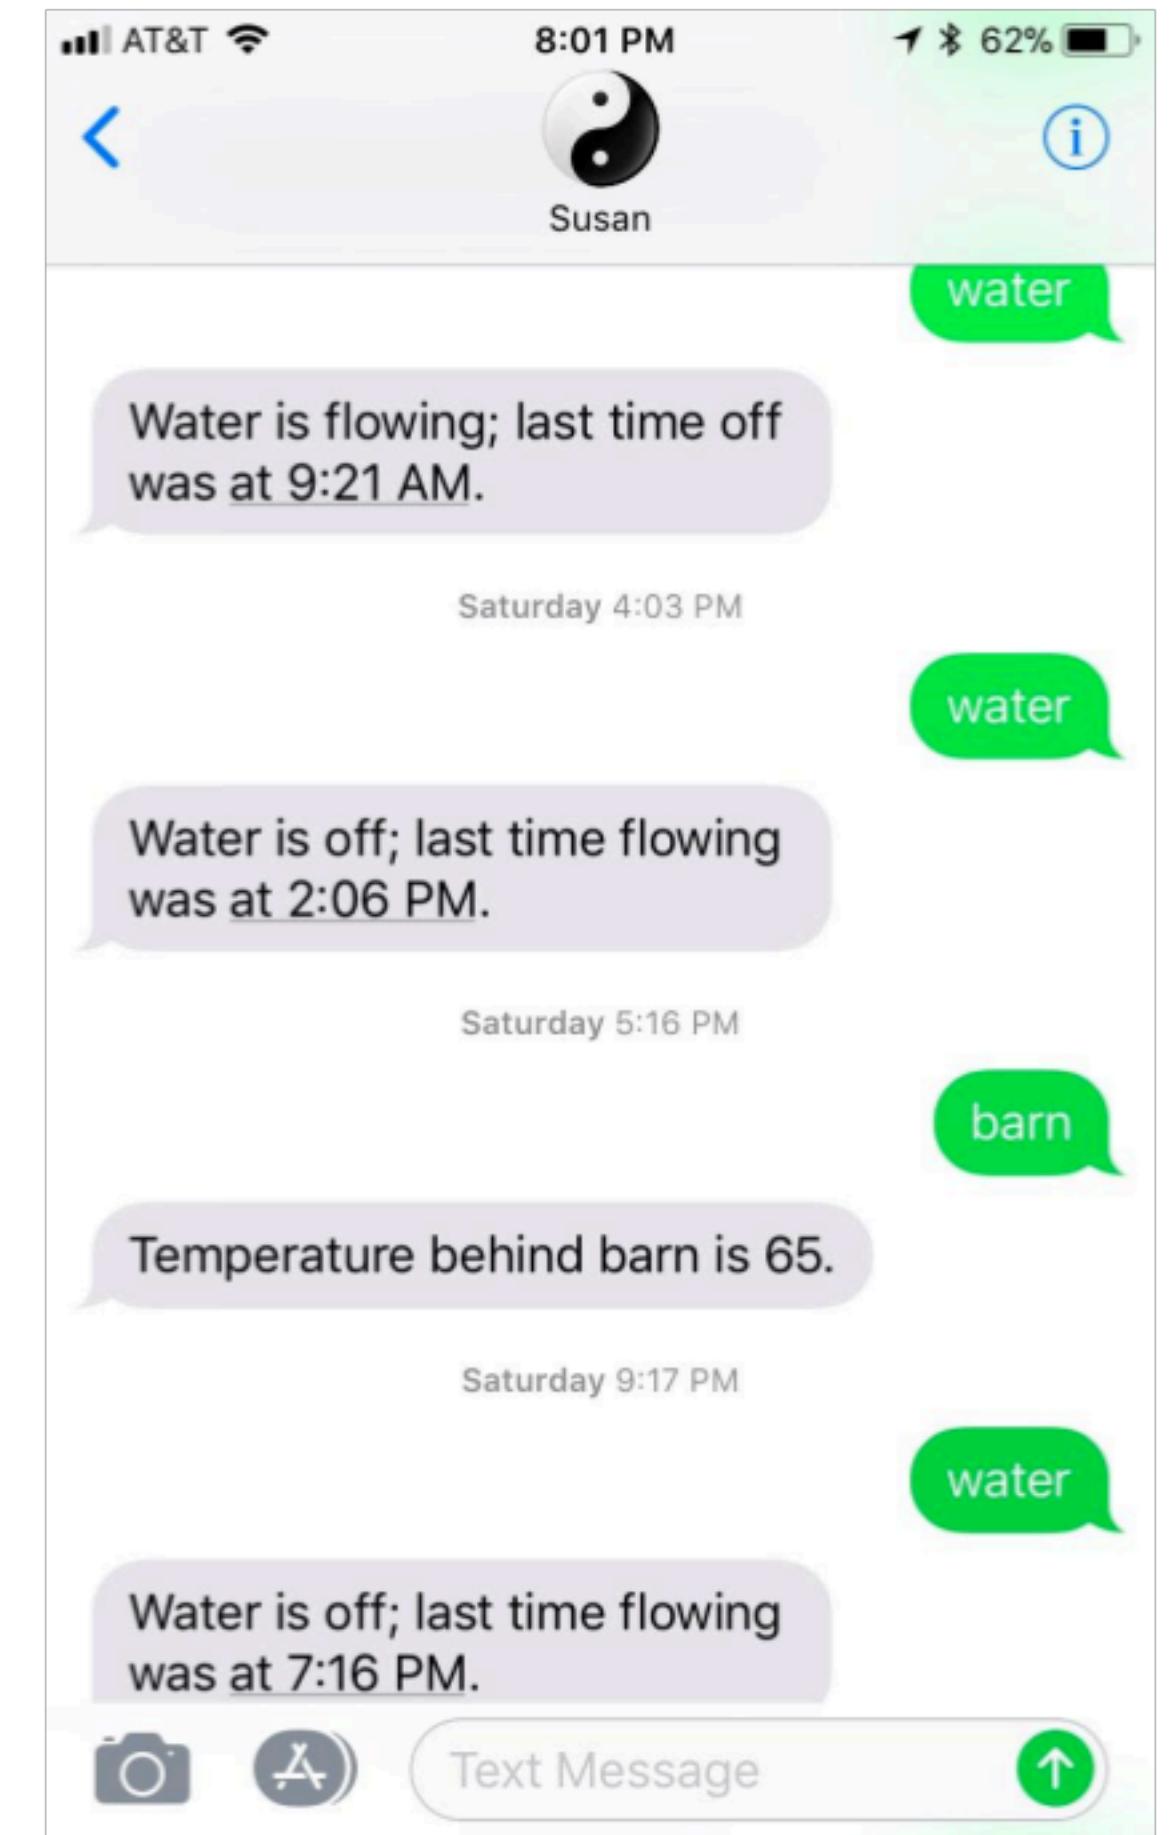
- Fast, memory efficient
- No central server or broker process needed
- ZMQ makes a great concurrency manager for multiple RPis (using synchronous REQ / REP)
- Communications protocol is flexible, but messaging protocol must be well designed
- Alternatives considered: RabbitMQ, ROS, AMQP, MQTT, others

Communications: Some Learnings

- Modeled on FAA protocol between ATC and pilots (pilot geekiness)
 - Who you are, where you are, what you want; predictable pattern
 - No unnecessary chatter; keep channel clear
 - Backup protocol if ATC is unavailable
- Use multiple communications channels
 - Transfer via ZMQ for some things: images and messages from RPis
 - Use scp & rsync for others: moving images around; backing things up
 - Dropbox and Google Drive are helpful in multiple ways

Communicating to End Users (me ;-)

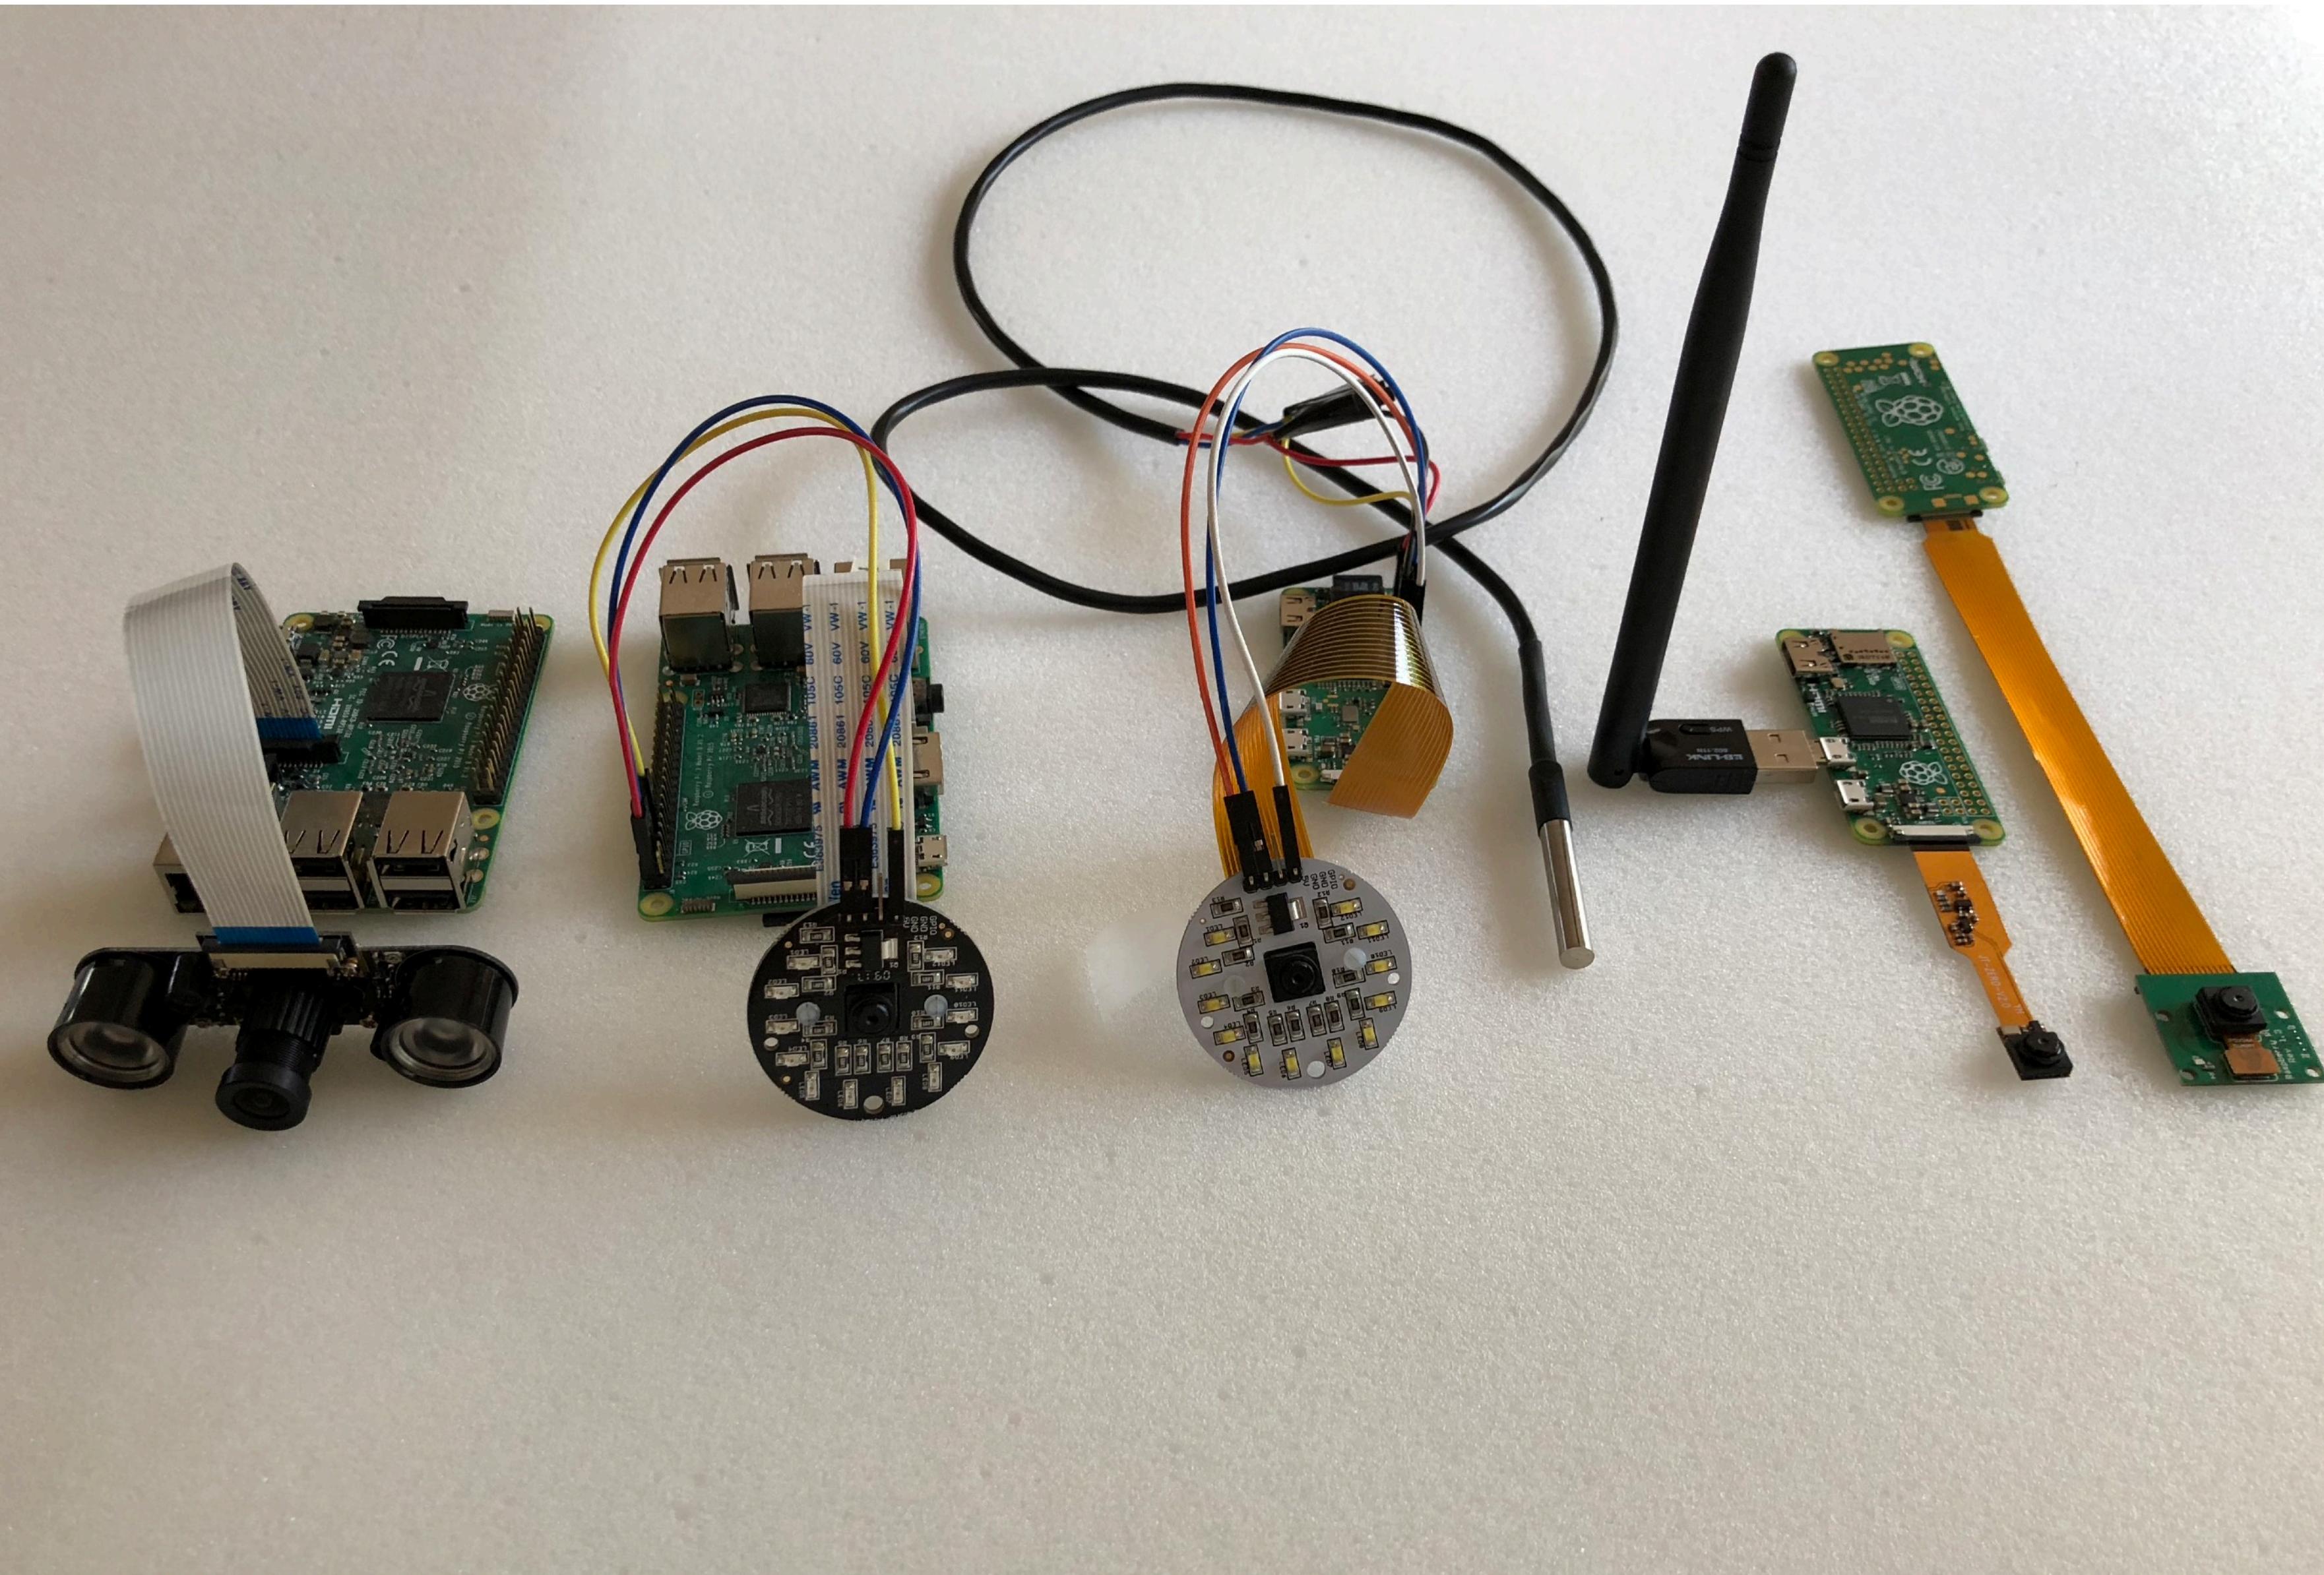
- Communicating “conclusions” from computer vision observations:
 - “Water is off; last time flowing at 7:15am”
 - “Garage is closed; last open at 8:20pm yesterday”
 - “Animal spotted at 3:15am; coyote (77%)”
- Communicating via multiple channels
 - Simple CLI text chat interface for testing
 - SMS texting interface has proved most useful so far
 - Website “dashboard” for images and data charts



Learnings: Hardware Tips and Tricks

- PiCams vs. WebCams: setting exposure mode, iso, contrast, brightness and saturation adjustment
- Enclosures, e.g. water meter cam, critter cam: mason jars and old shingles are my DIY favorites
- Cost of different cams: the variation is large; the variation in quality is even larger
- InfraRed lighting for critters: multiple choices for size and on/off control, including PIR sensors
- LED lighting under RPi control: MOSFETs rule!
- Temperature / humidity sensors: DS18B20 & W1ThermSensor; DHT &
- Networking via power-line modem ethernet as well as WiFi and wired ethernet
- Laptops keep working during power failures; RPis don't; design considerations
- Power: advantages of 12v over 5v: cheap reliable converters (yes, those car chargers for your phone)

Hardware: Lights, Camera, Action!



Left to Right:

RPi with Waveshare
Combo IR lens and dual
IR floodlights

PiNoir IR Camera with
IR “Ring Light” floodlight

RPi Zero with PiCamera
in white light “Ring Light”
with DS18B20 temperature
probe

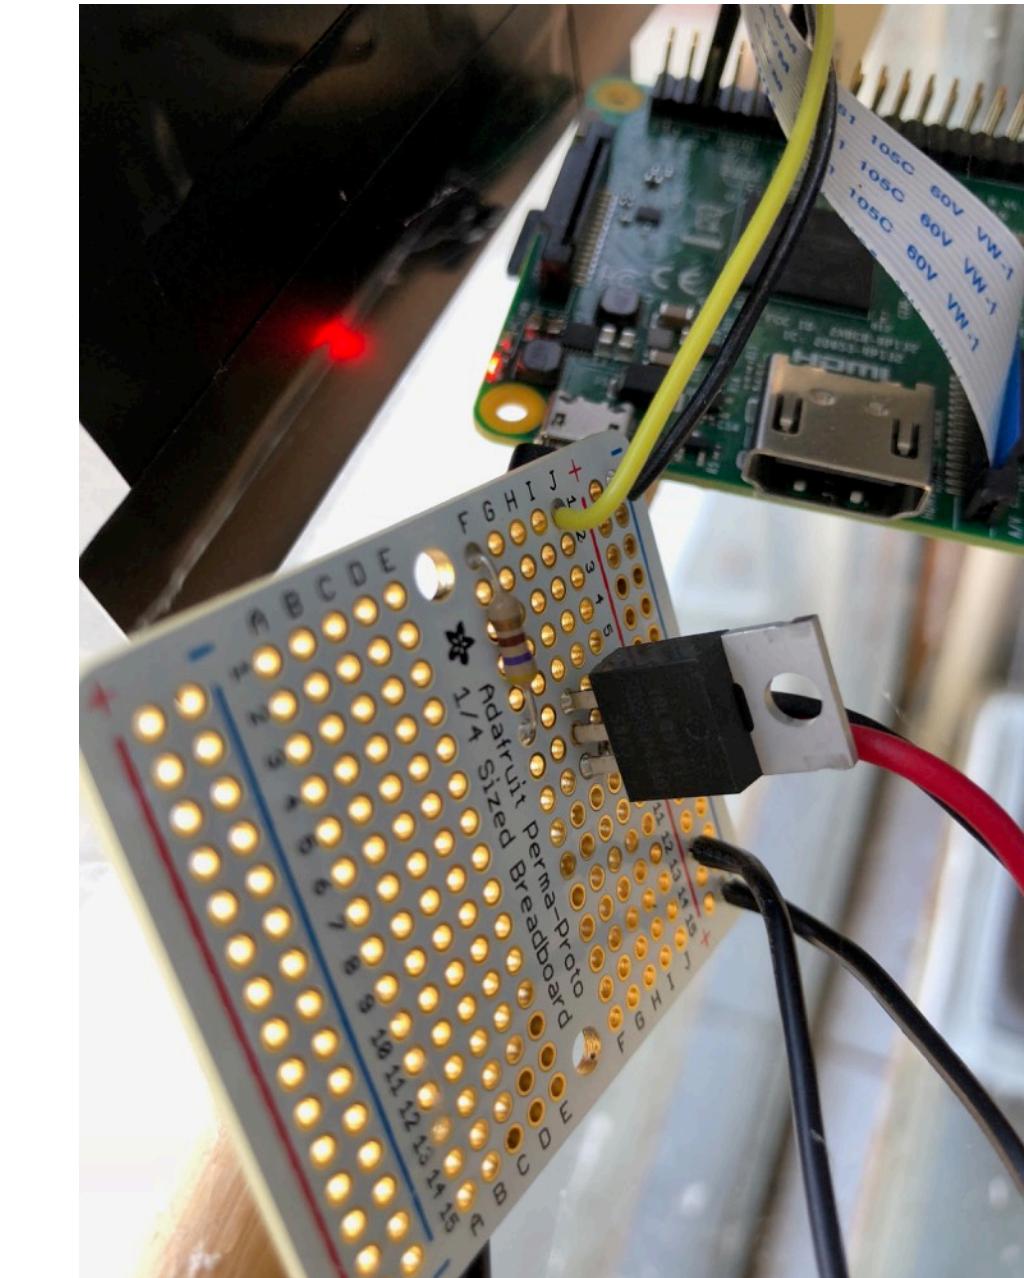
RPi Zero with “Spy Cam”
and longer range WiFi

RPi Zero with older model
(and half price) PiCamera

Hardware: More Light Types

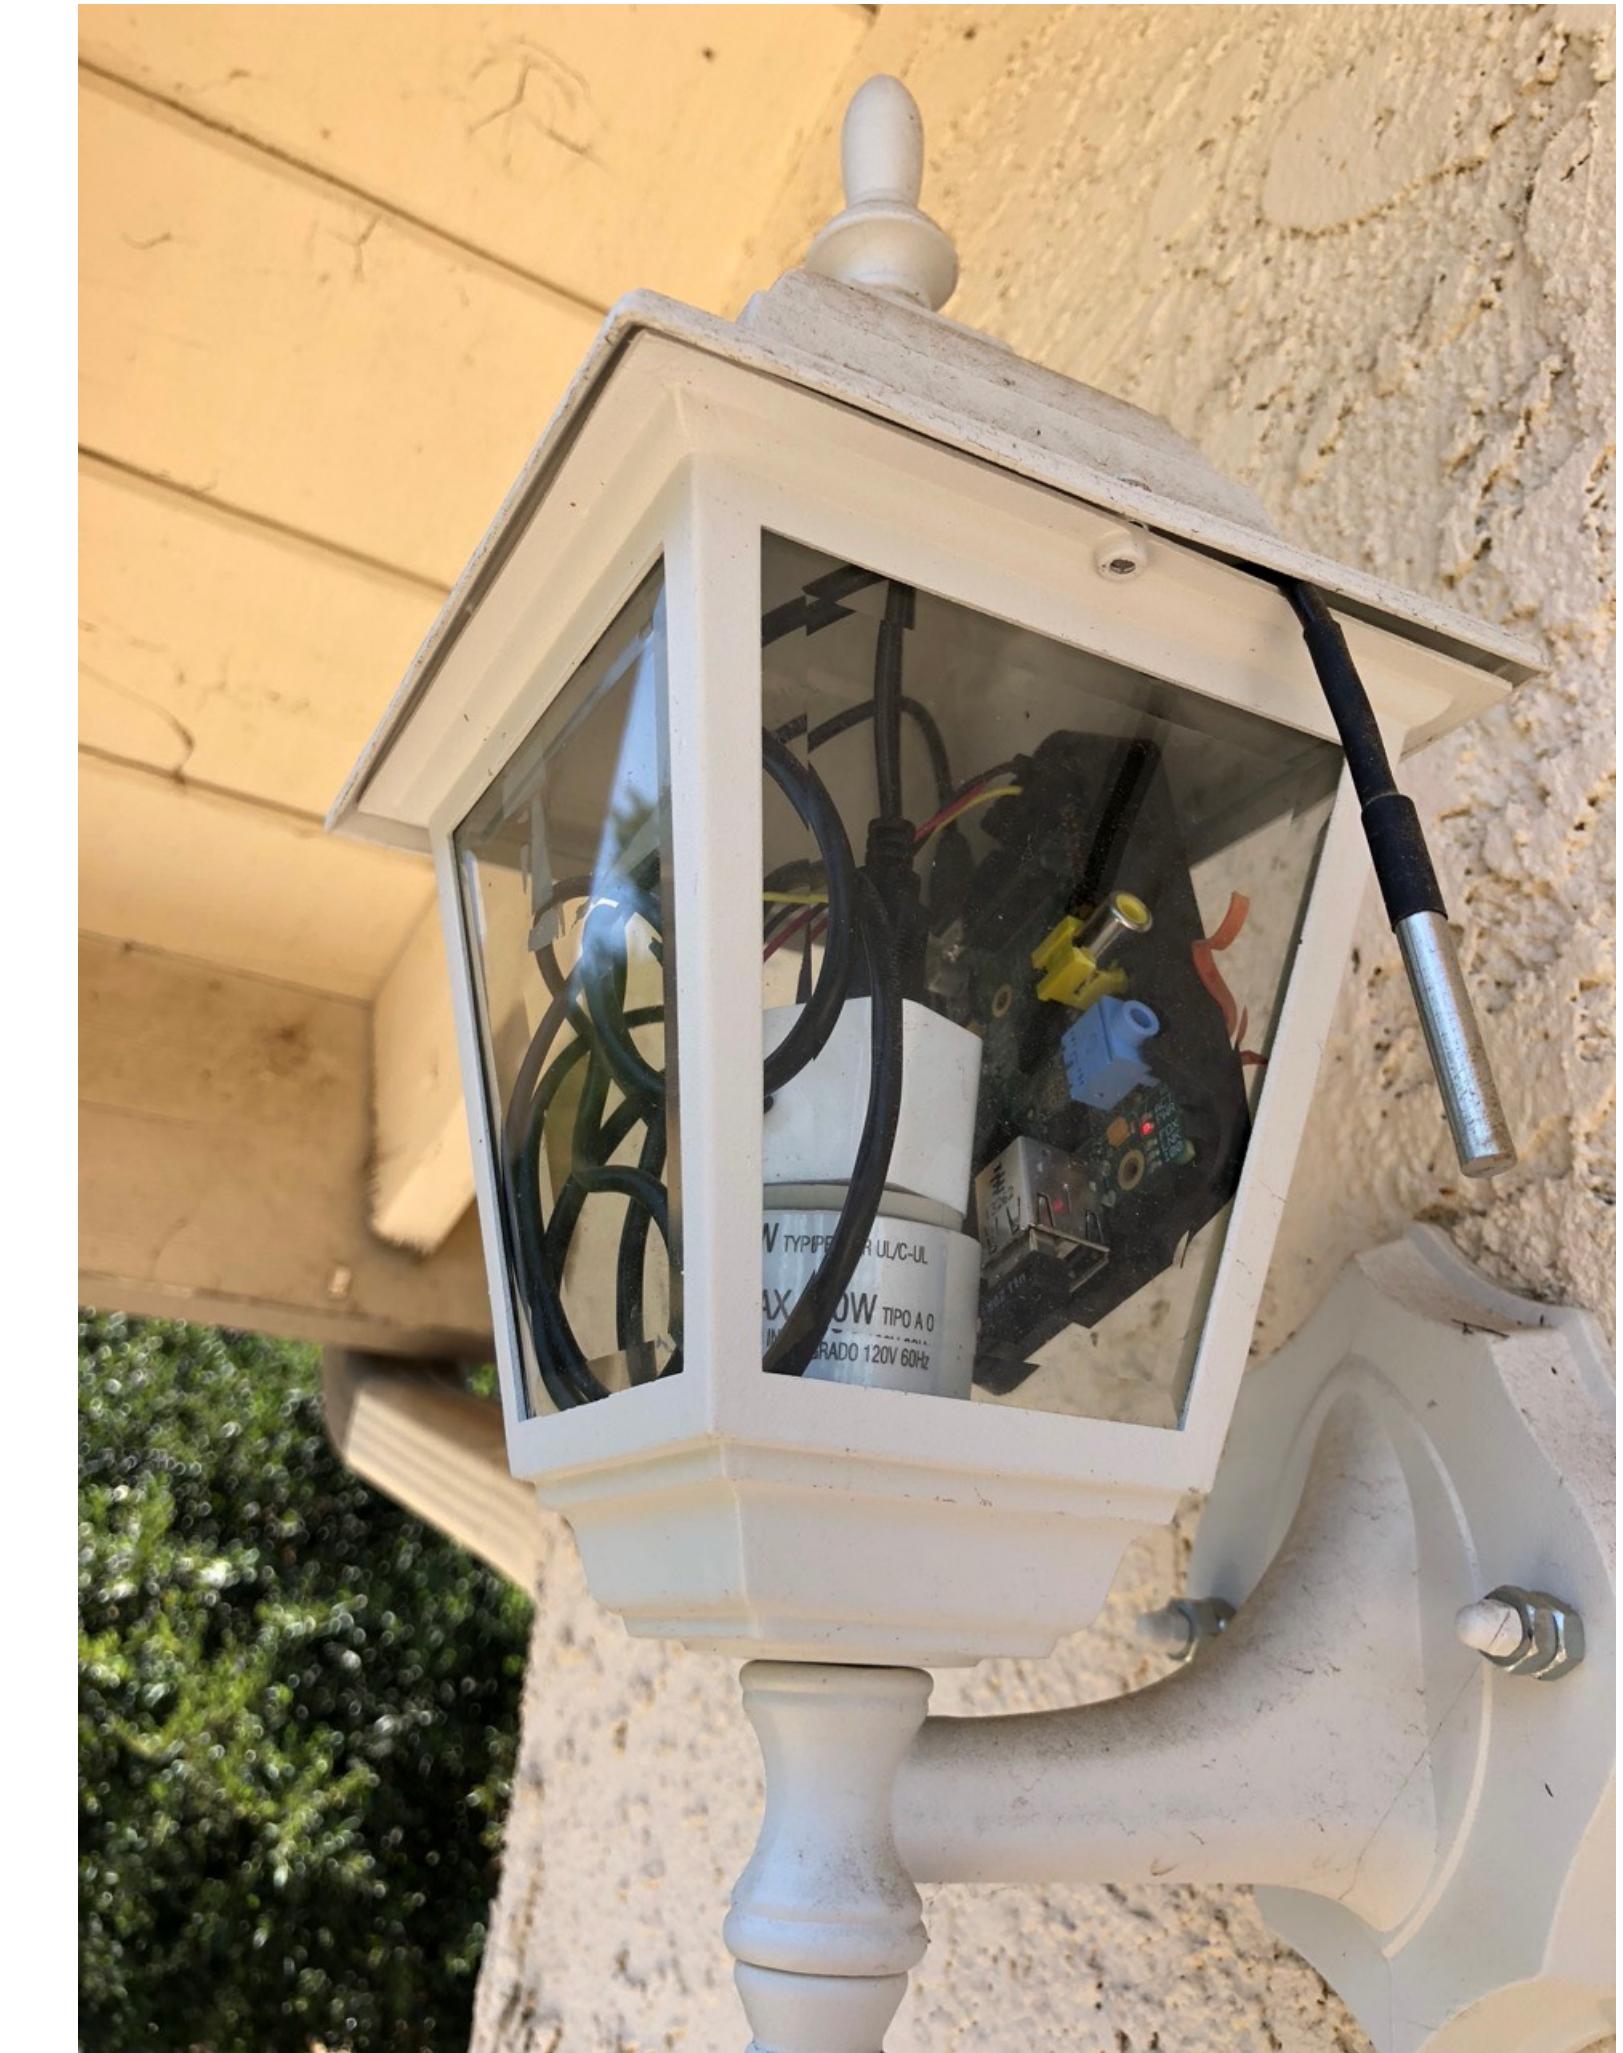
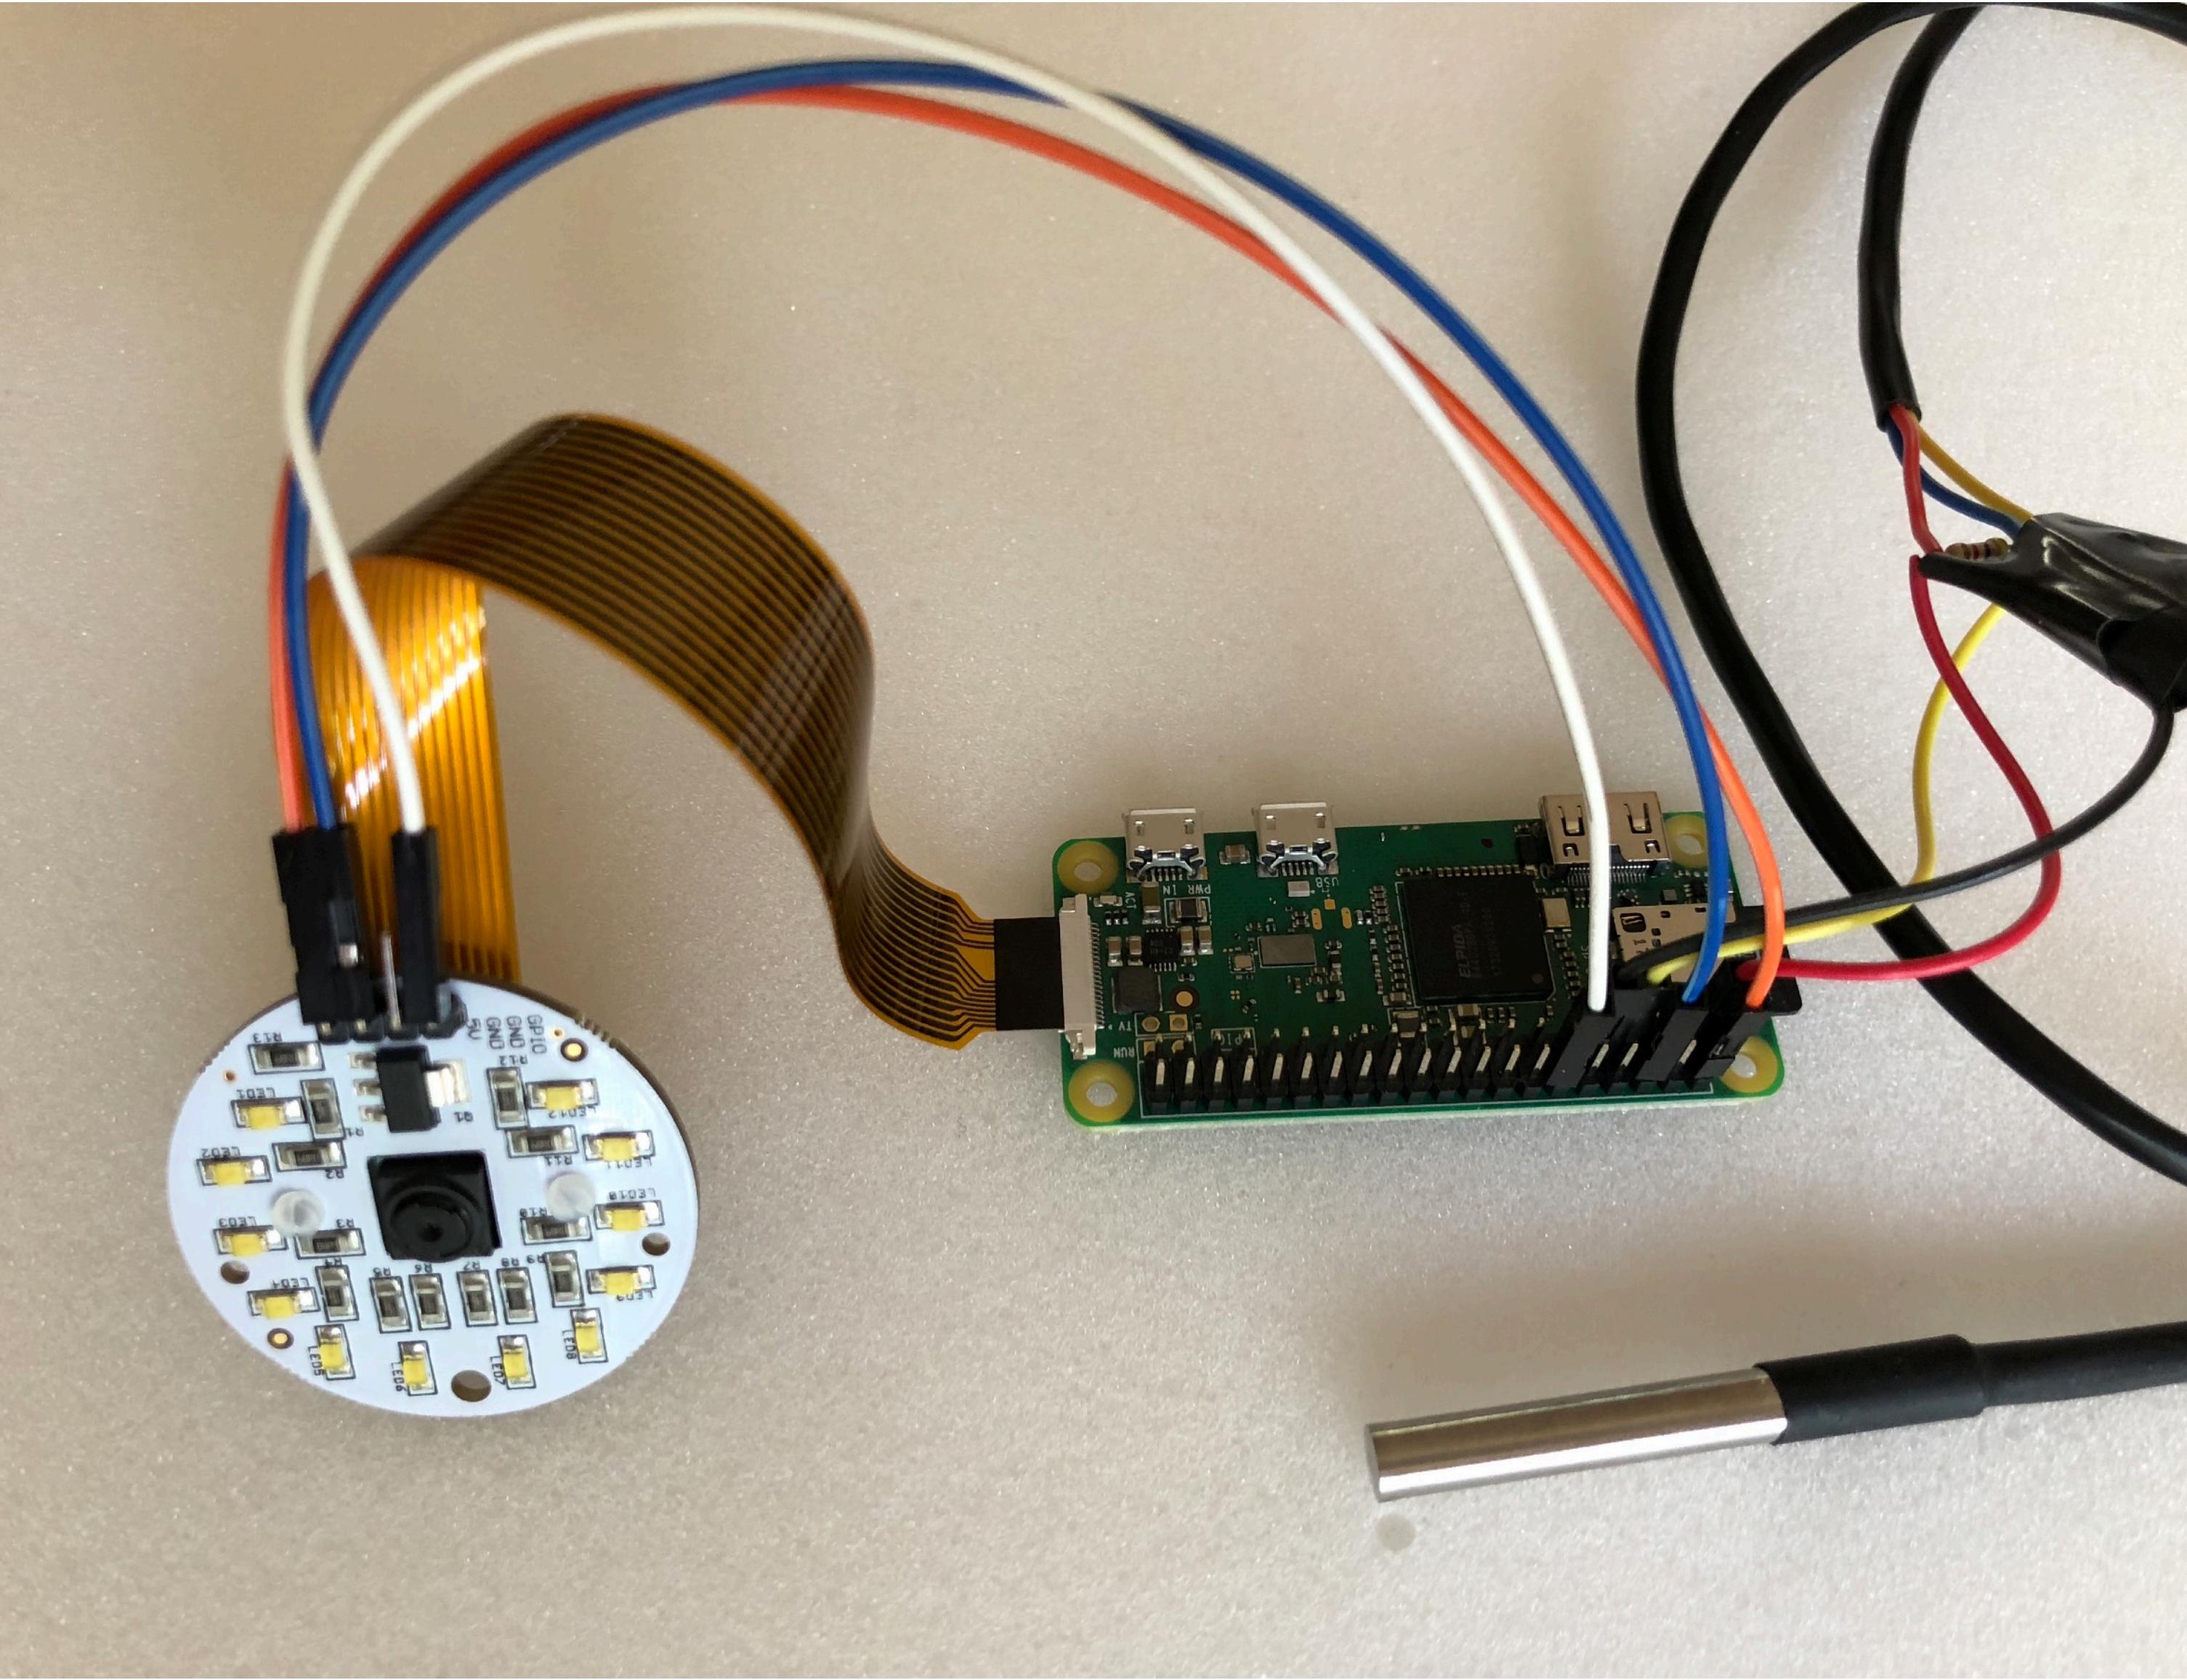


White floodlight (18W), 12V power cube, Infrared floodlight (8W)



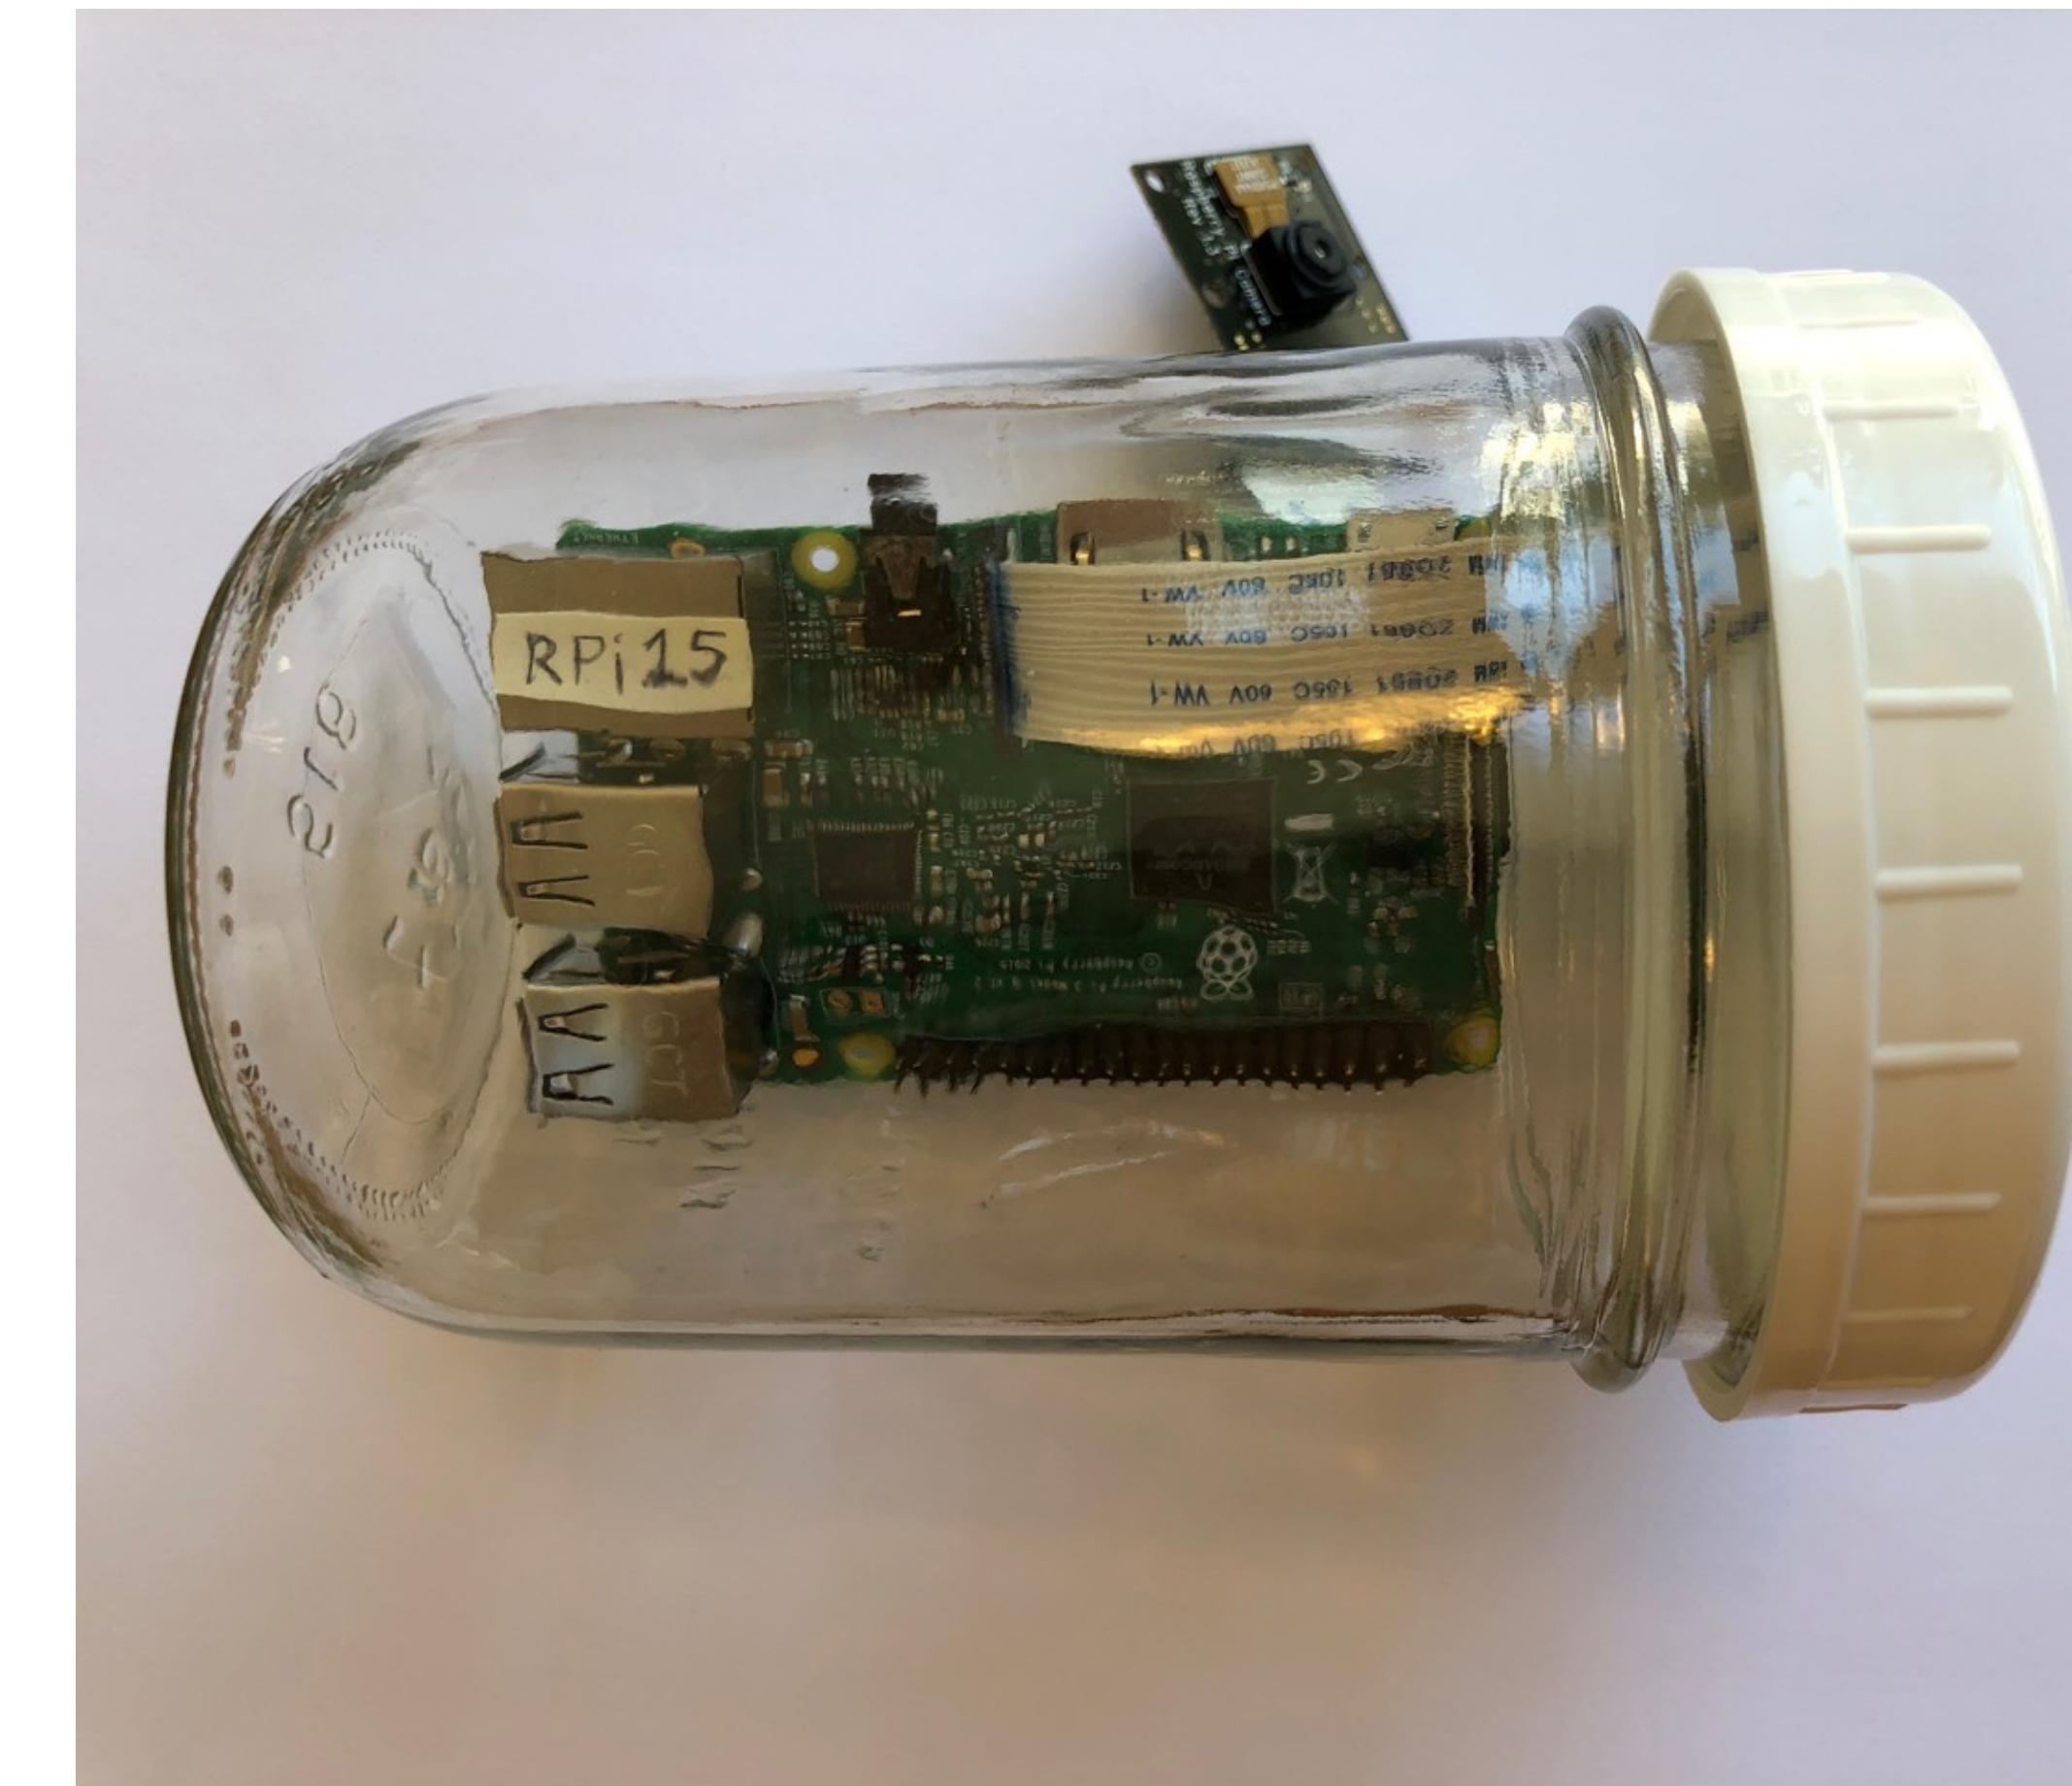
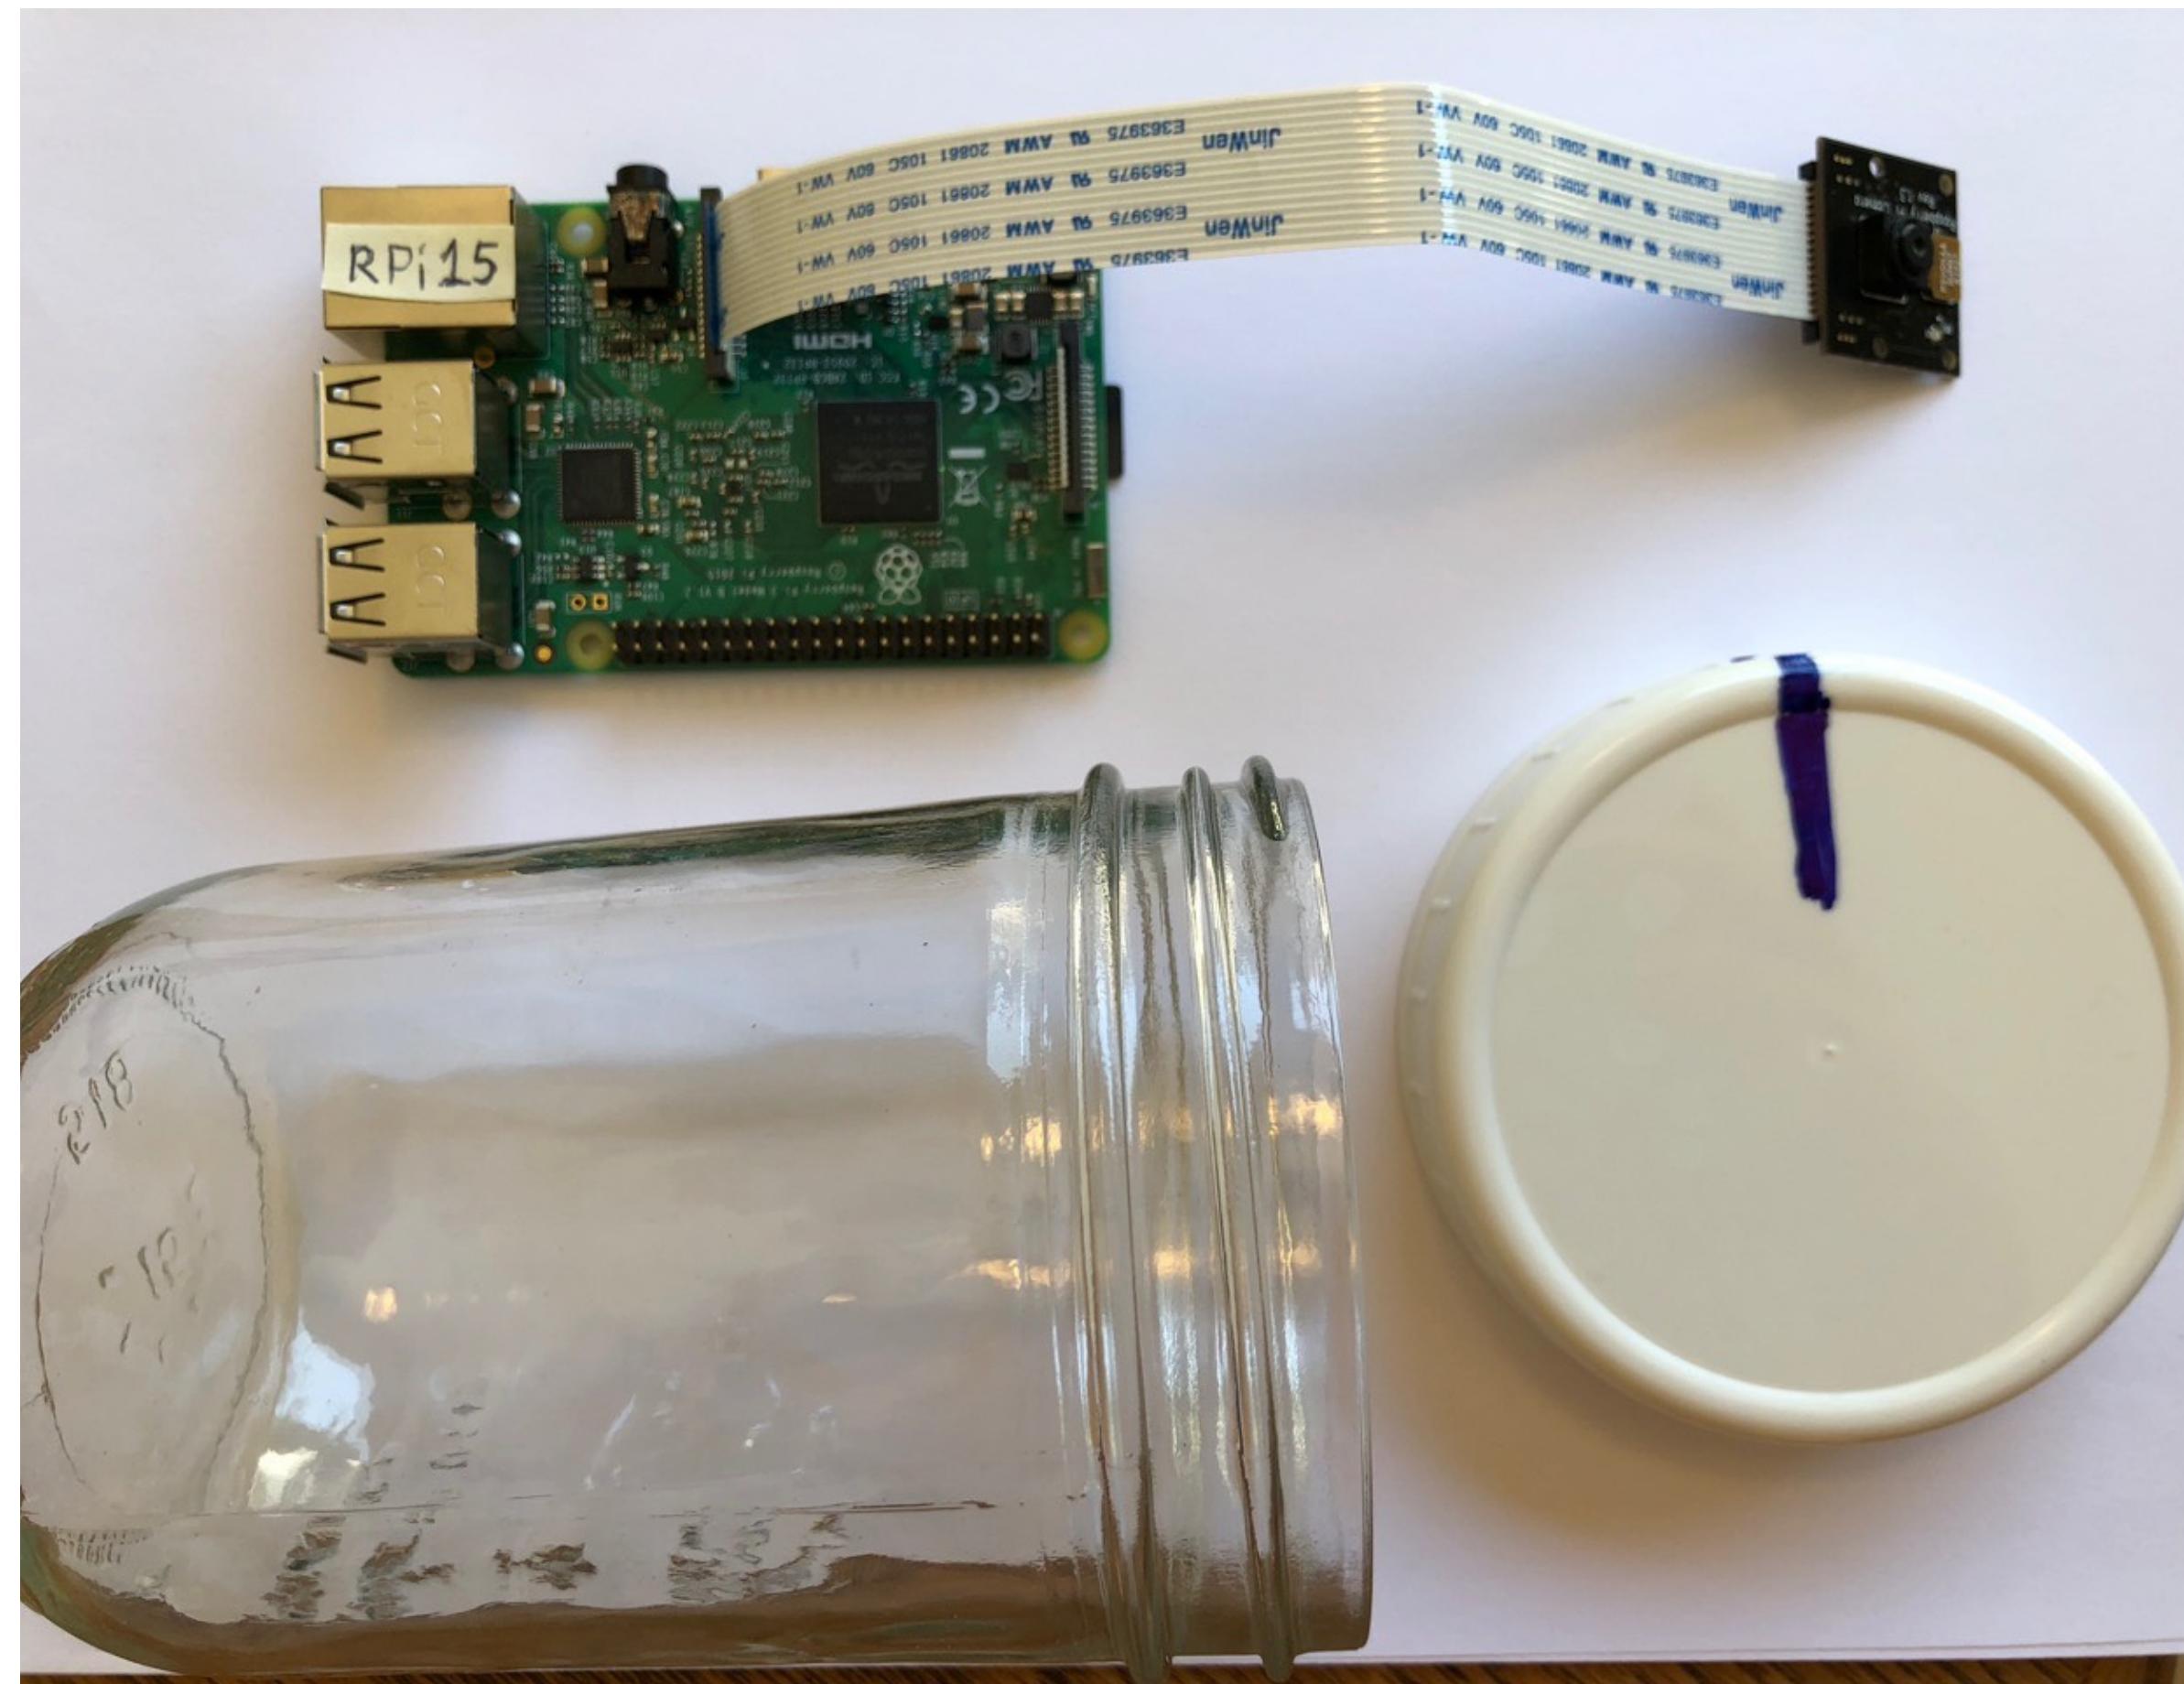
Simple MOSFET circuit
provides RPi GPIO
control for floodlights

Hardware Enclosures: Light Fixtures!

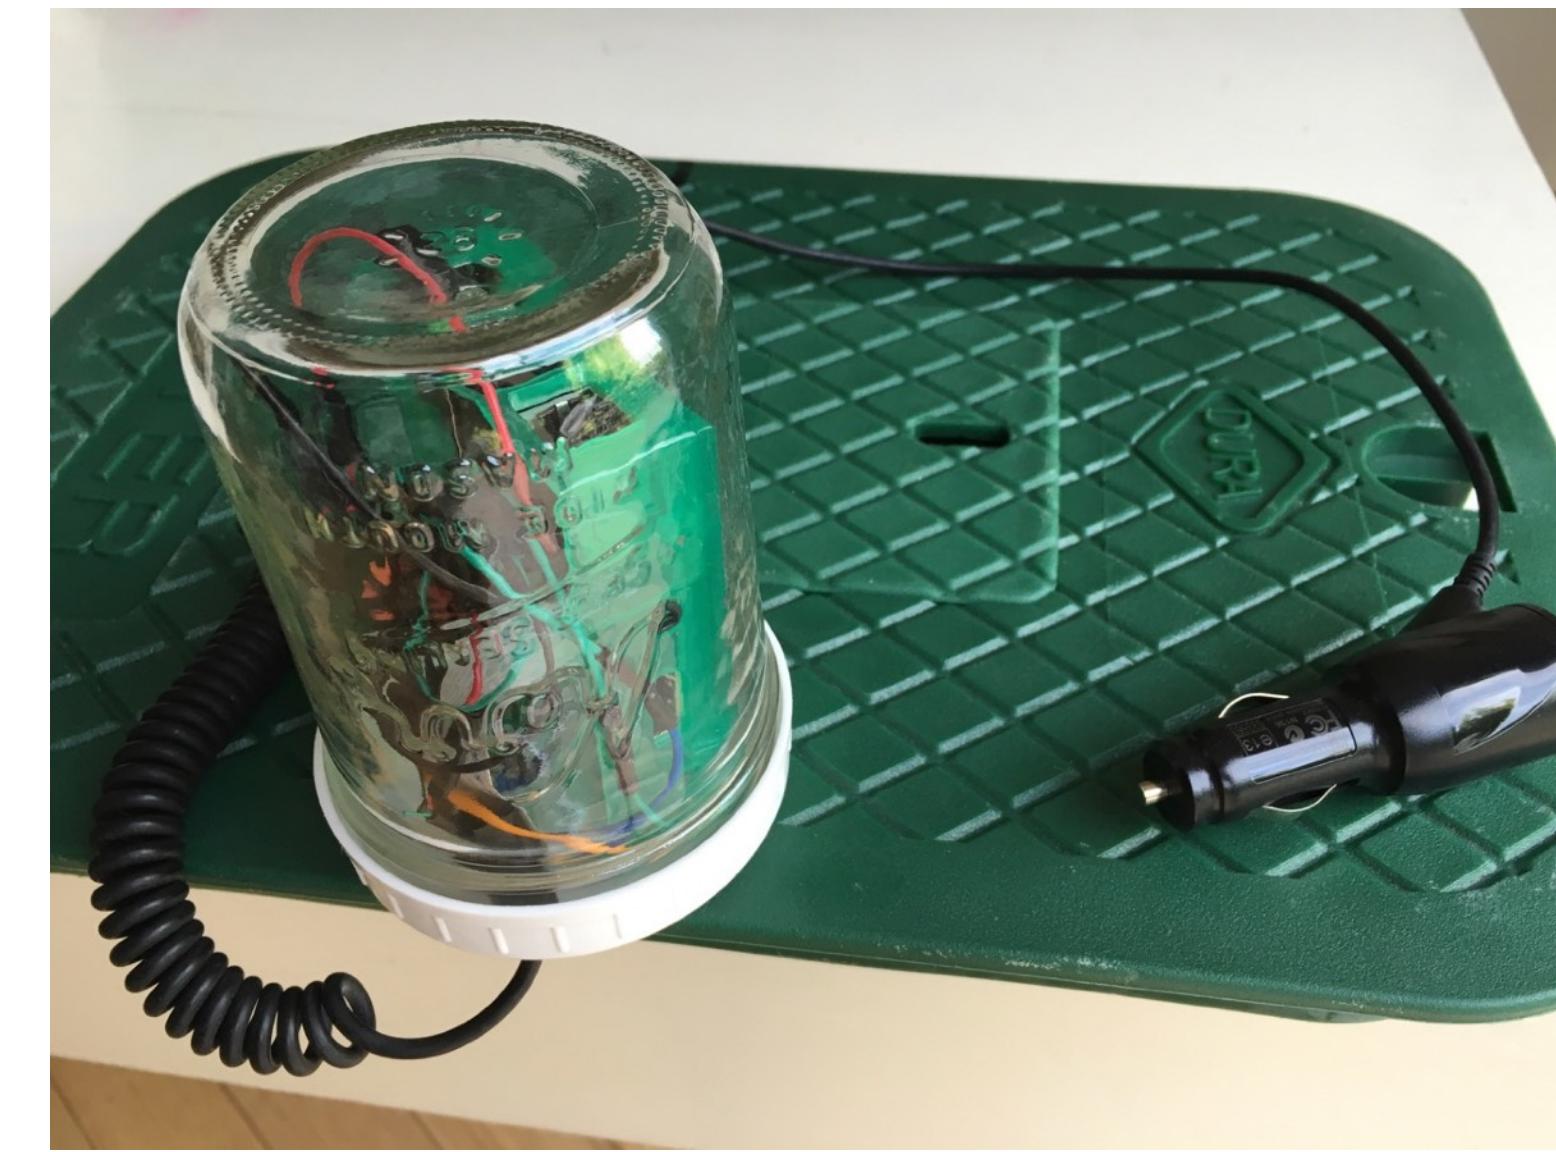
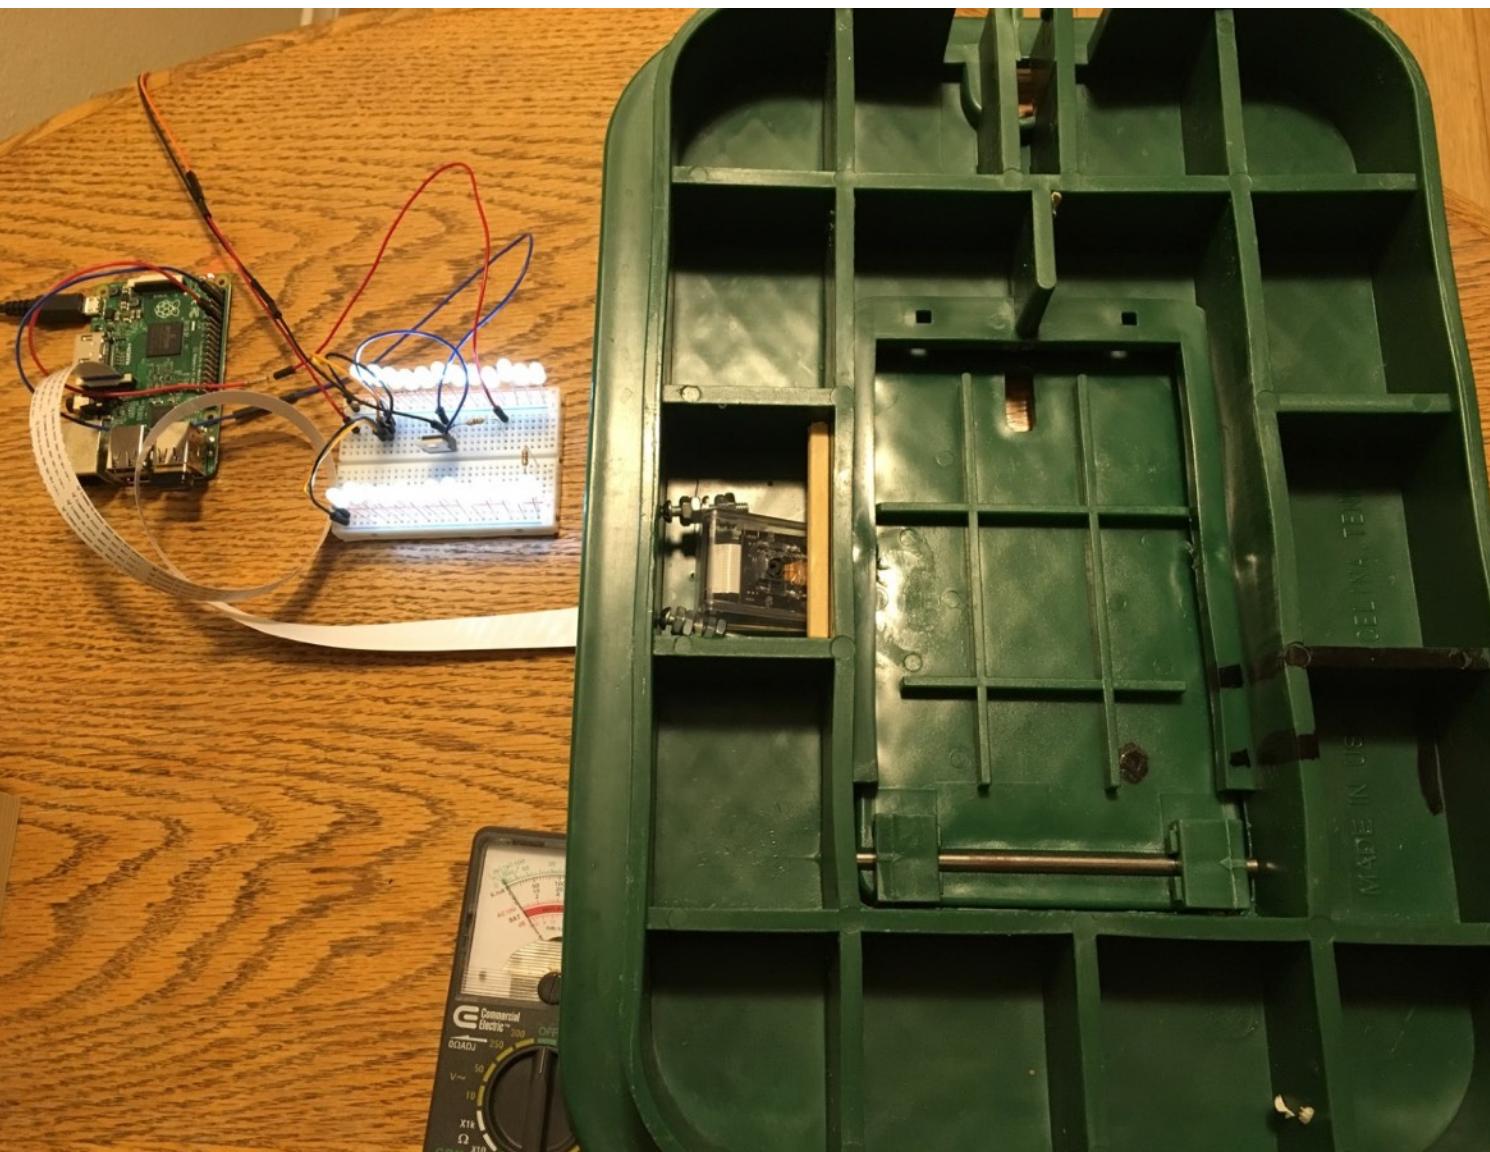
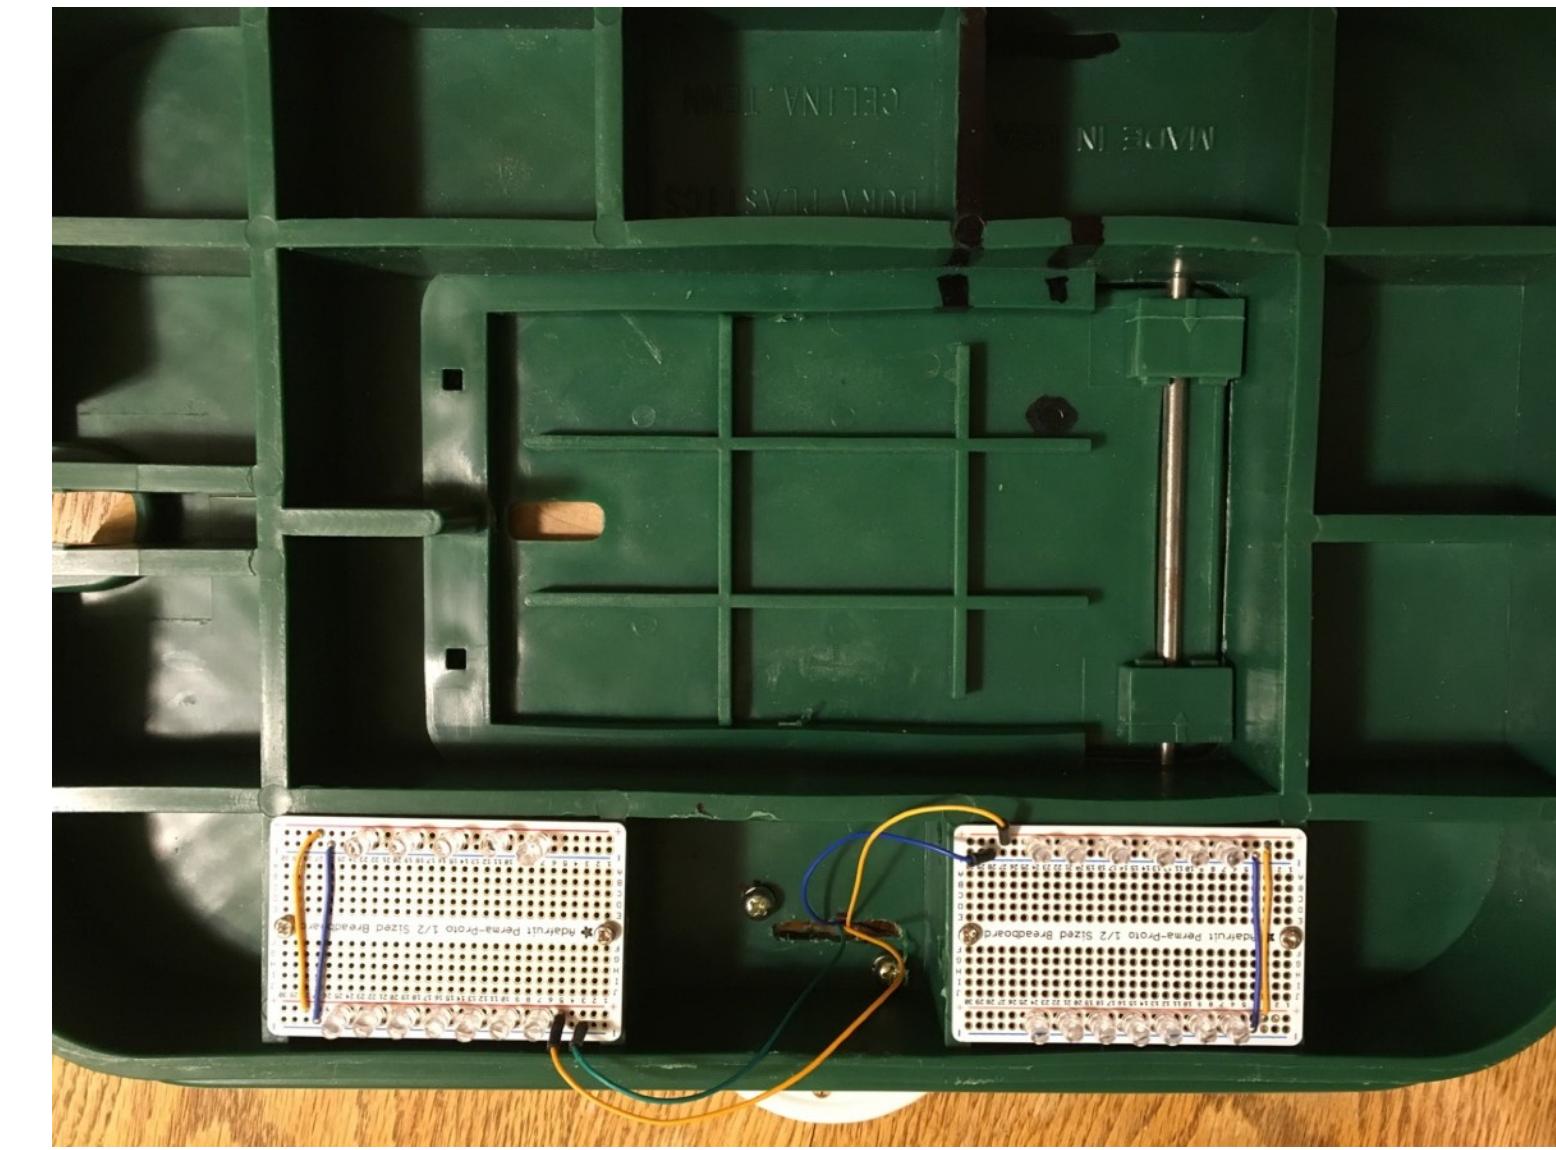
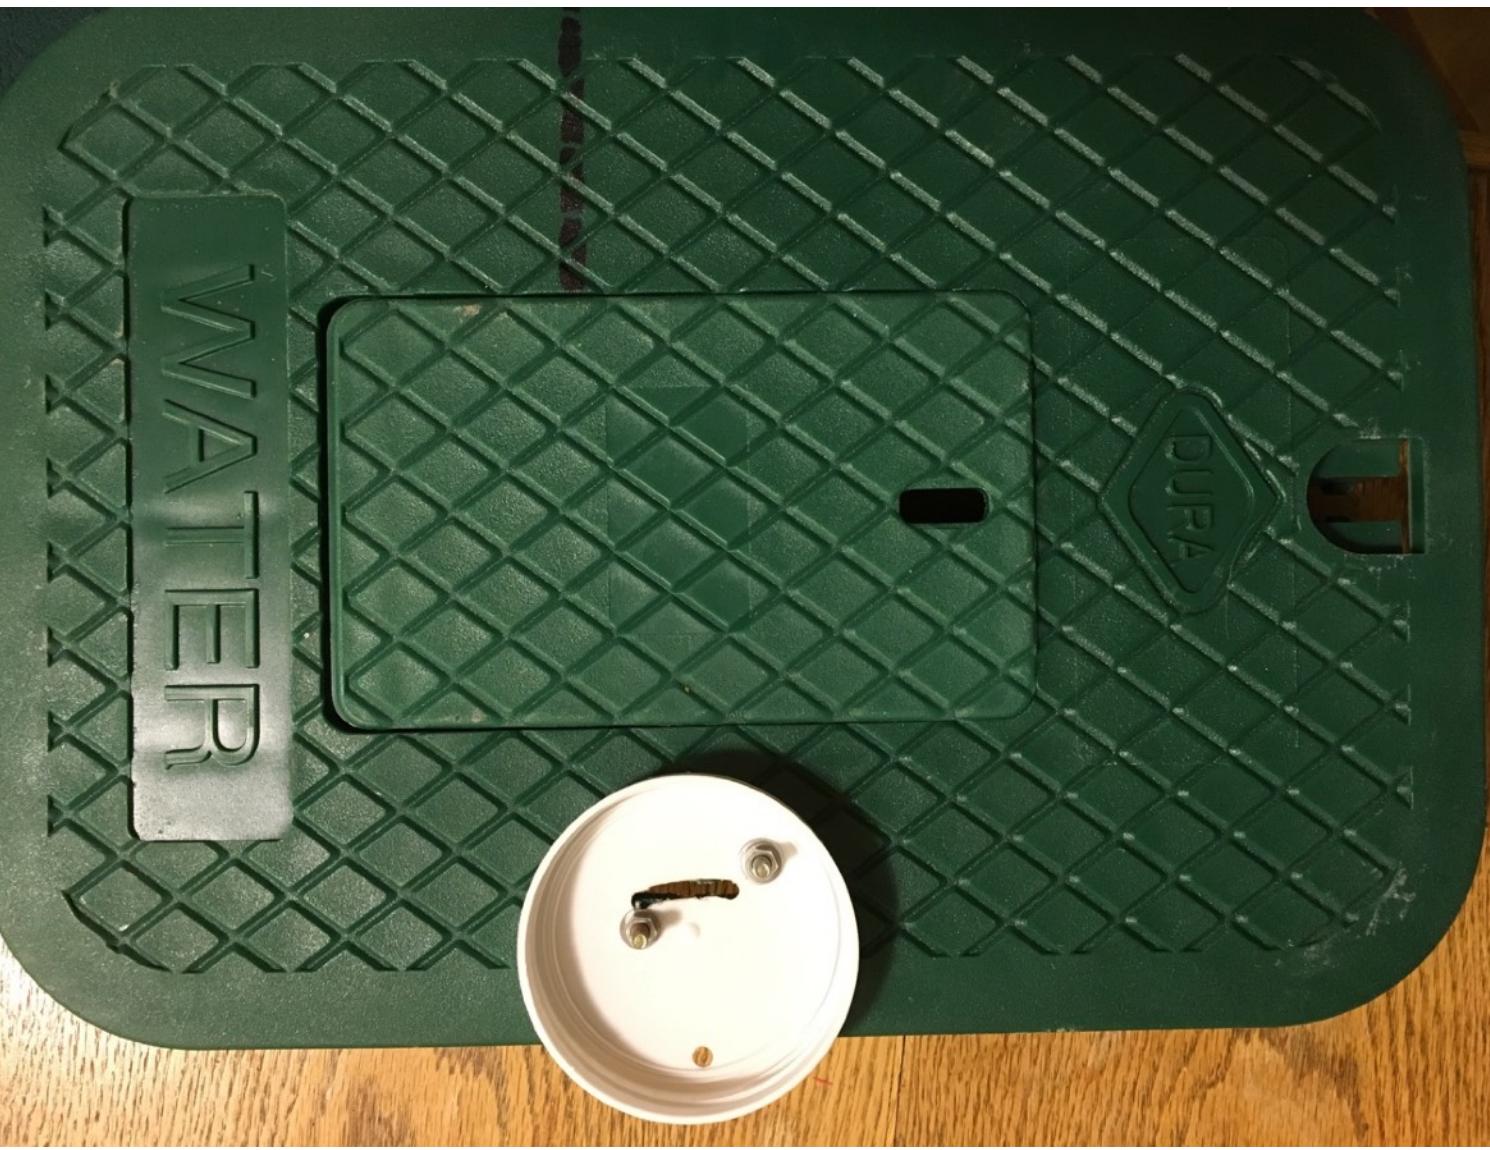


... and temperature sensors!

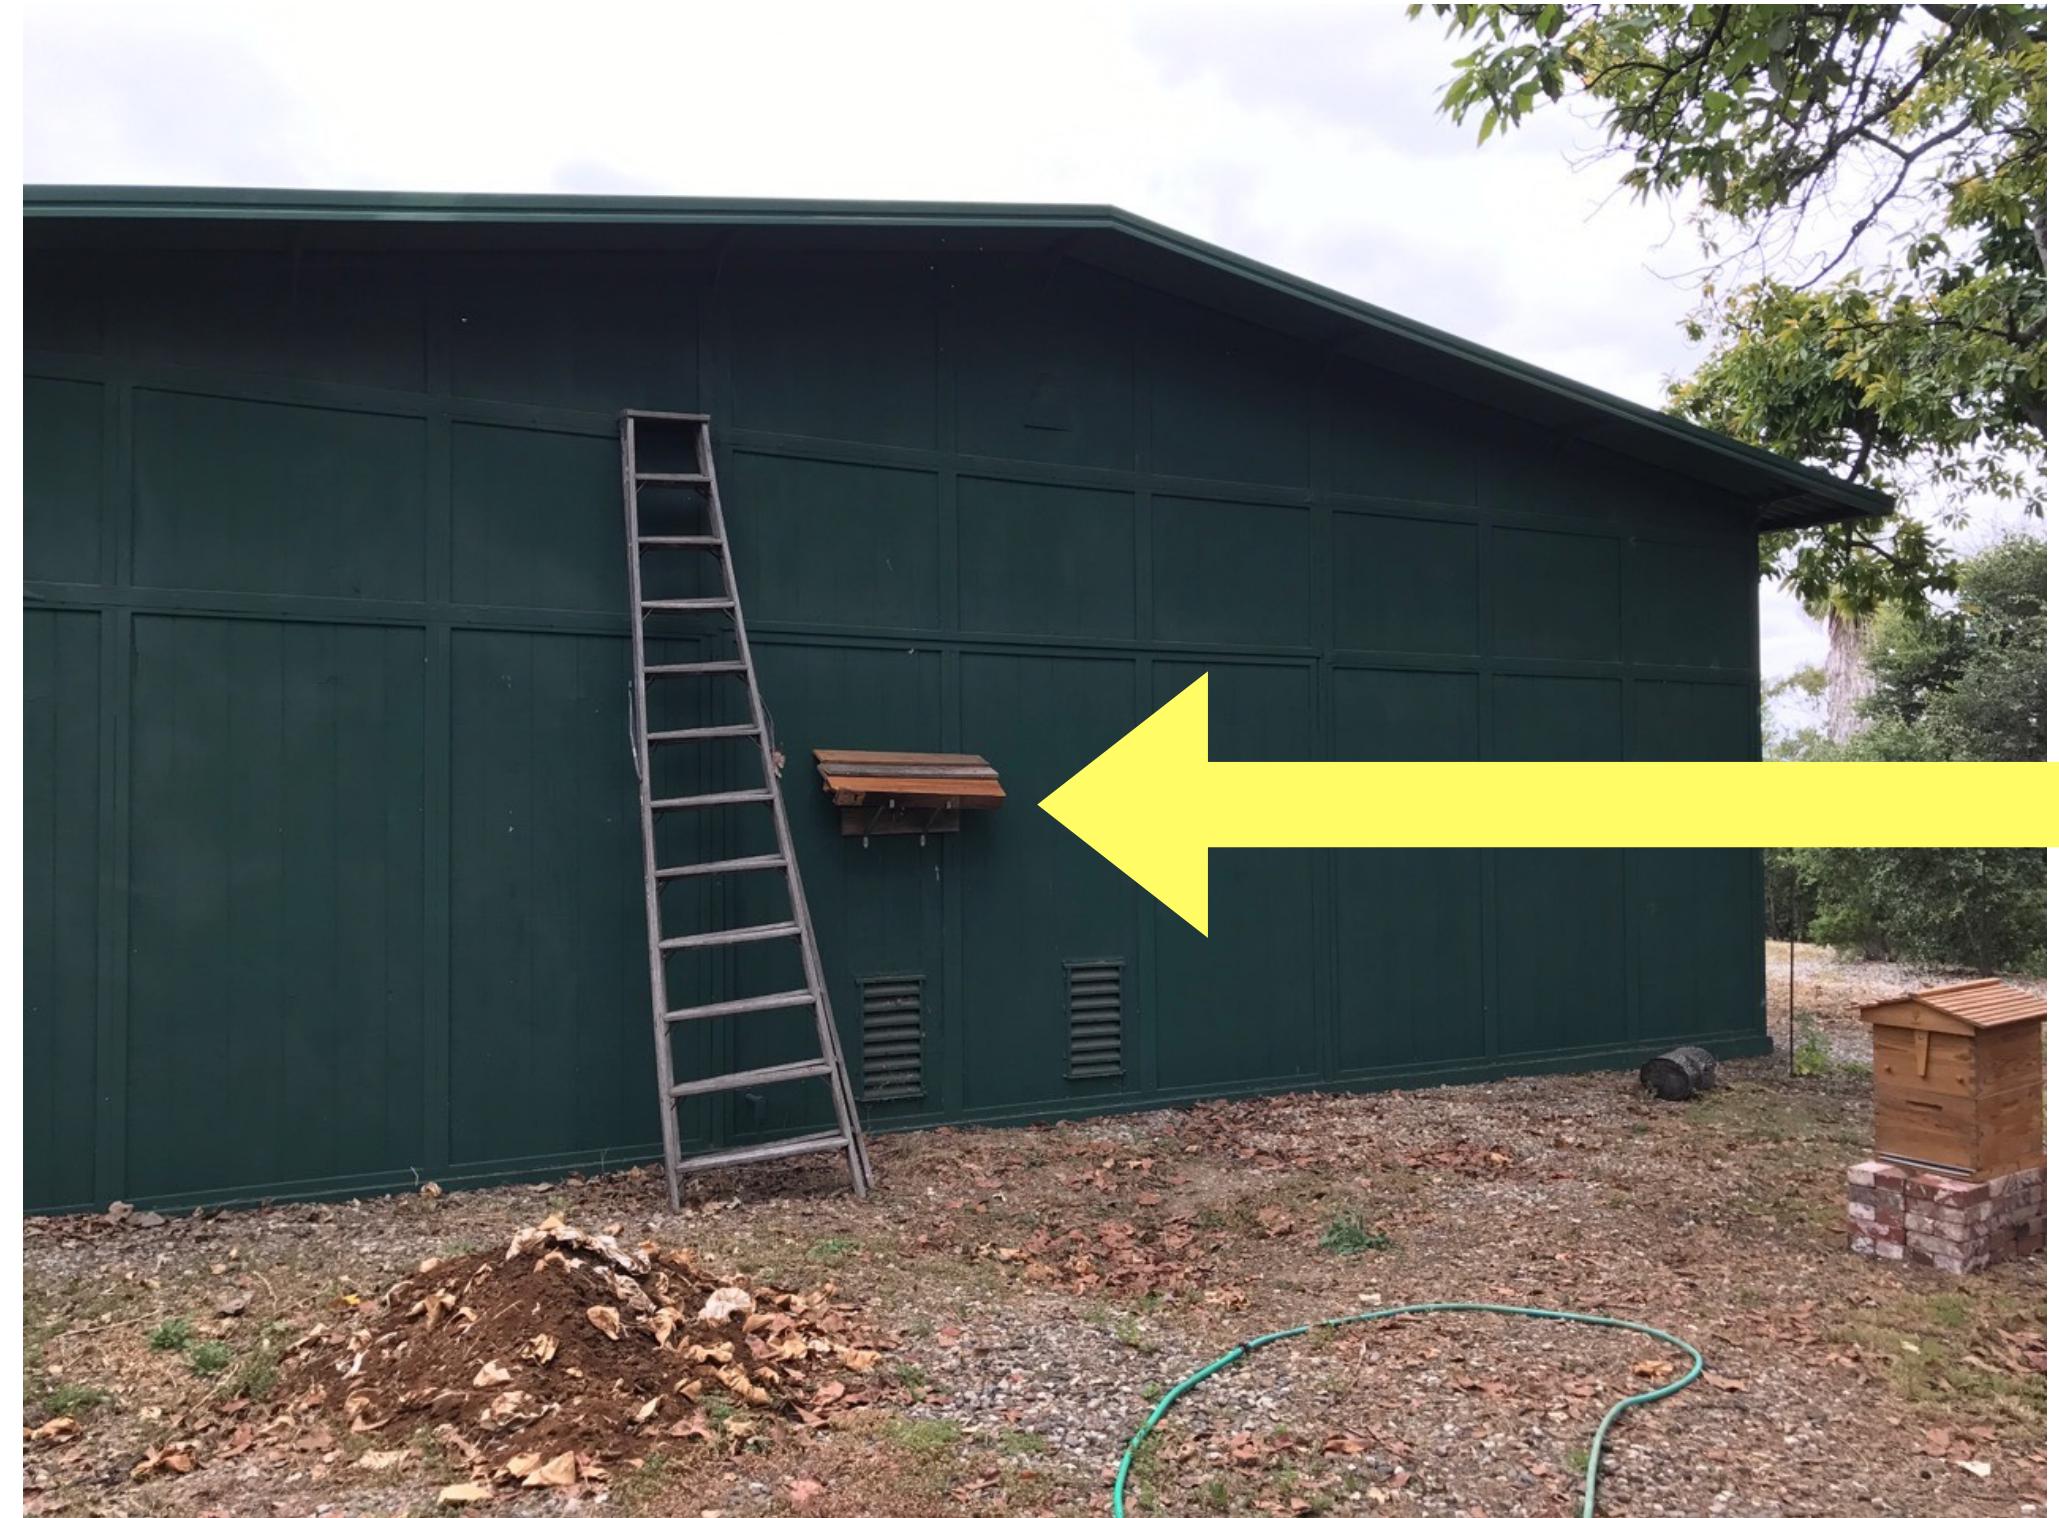
Hardware Enclosures: Mason Jars!



Hardware Enclosures: Water Meter Cam



Hardware Enclosures: Barn Coyote Cam



**Glass blocks Infrared, so Mason jars
can't be used for Infrared Cameras.
But a PiNoir camera does fine outdoors
under a couple of old cedar shingles.**



Interlude: Coyote Cameos with Rabbit



**Taken with RPi
“PiNoir” Infrared camera,
using an 8W Infrared
floodlight**

Hardware Enclosures: Fake Security Cam for \$5



This \$5 fake security cam can be modified to hold a RPi Zero with PiCamera.

It is waterproof and has tripod quality adjustable view angle.

A plain RPi zero case is usually \$5.



Some Lessons Learned So Far

- Experiment & observe before optimizing; change, document carefully & iterate
- Tune computer vision parameters carefully (thresholds, pixel area parameters, etc.)
- Get the lighting right! Control the lighting brightness with RPi (PWM & MOSFETs)
- Optimize network load; use isolated subnetworks (subnetworks add security, also)
- Reliability planning: recovery from power failures, internet outages, etc.
 - What to prioritize for battery backup; use laptops (have batteries) as hubs
 - When to connect with WiFi vs. wired ethernet
 - Self-restarting subsystems (init.d scripts versus systemd services)

Next: Indicator Plants to Measure Deep Soil Moisture



The most effective soil moisture test is an “indicator” plant that has deep roots and wilts before the plants around it do. Observation of such an “indicator” plant is a great use case for computer vision.

Interlude: Using project like this as a Resume Builder

- Build a portfolio that shows **broad practical experience**; even “hobby” projects are helpful
- Document your portfolio; emphasize specific examples and learnings (what’s special about it?)
- Demonstrate ability to **complete** large projects fully (80% solutions are good; 100% solutions show ability to finish)
- Show attention to project details (and to project **management** details)
- Examples of iteration and continuous improvement, including iteration of design
- **Effective documentation** for yourself, for others, for long term maintenance; for publications
- Effective use of Git/GitHub (or other version control) as a collaboration tool
- Communication skills (including to non-tech audience) and team collaboration skills
- Be able to tell a **short, compelling project story**; the “elevator speech” about your big project

Next Steps for This Project

- Deep learning for more specific conclusions from these computer vision pipelines
 - Reliably distinguish coyote vs raccoon; butterfly vs hummingbird; census counts
 - Recognize and label specific critter individuals (I can tell 2 different coyotes apart from a video; will develop code to have computer vision / machine learning do it)
 - Use CV / DL to classify “indicator plants” as “water stressed” (or not) in real time
- Measure partly cloudy vs hazy sun; compute effective value of sunlight for cumulative plant photosynthetic growth; visually measure water stress of plants
- Natural Language Processing for better communication of computer vision results
- Fine tuning parameters via interactive text commands (rather than changing yaml files)...or... using deep learning to optimize parameters?
- Implementing IOT actions: close the garage; irrigate specific garden areas, etc.

Helpful Resources

- **Electronics Resources:**

- Adafruit has extremely helpful tutorials as well as a deep selection of electronic parts and kits: <https://www.adafruit.com/>
- Sparkfun also has great tutorials and electronic goodies; lots of cameras; lots of fun gadgets: <https://www.sparkfun.com/>
- **ZMQ:** <http://zeromq.org/> and the Python bindings: <https://pyzmq.readthedocs.io/en/latest/>

- **Permaculture:**

- Toby Hemenway, *Gaia's Garden A Guide to Home-Scale Permaculture*. Also by Toby: *The Permaculture City: Regenerative Design for Urban, Suburban and Town Resilience*
Link: <http://a.co/izt7Xun>
- **To learn more about (or contribute to!) this project:**

- <https://github.com/jeffbass>